



Training Days 2019 - February 19-21 - Westin Westminster, Colorado

---

# MongoDB 101 Workshop

## Sig Narváez





# Who Am I?

---



Sigfrido “Sig” Narváez

---

Principal Solution Architect  
MongoDB  
[@SigNarvaez](https://twitter.com/SigNarvaez)

Disclaimer



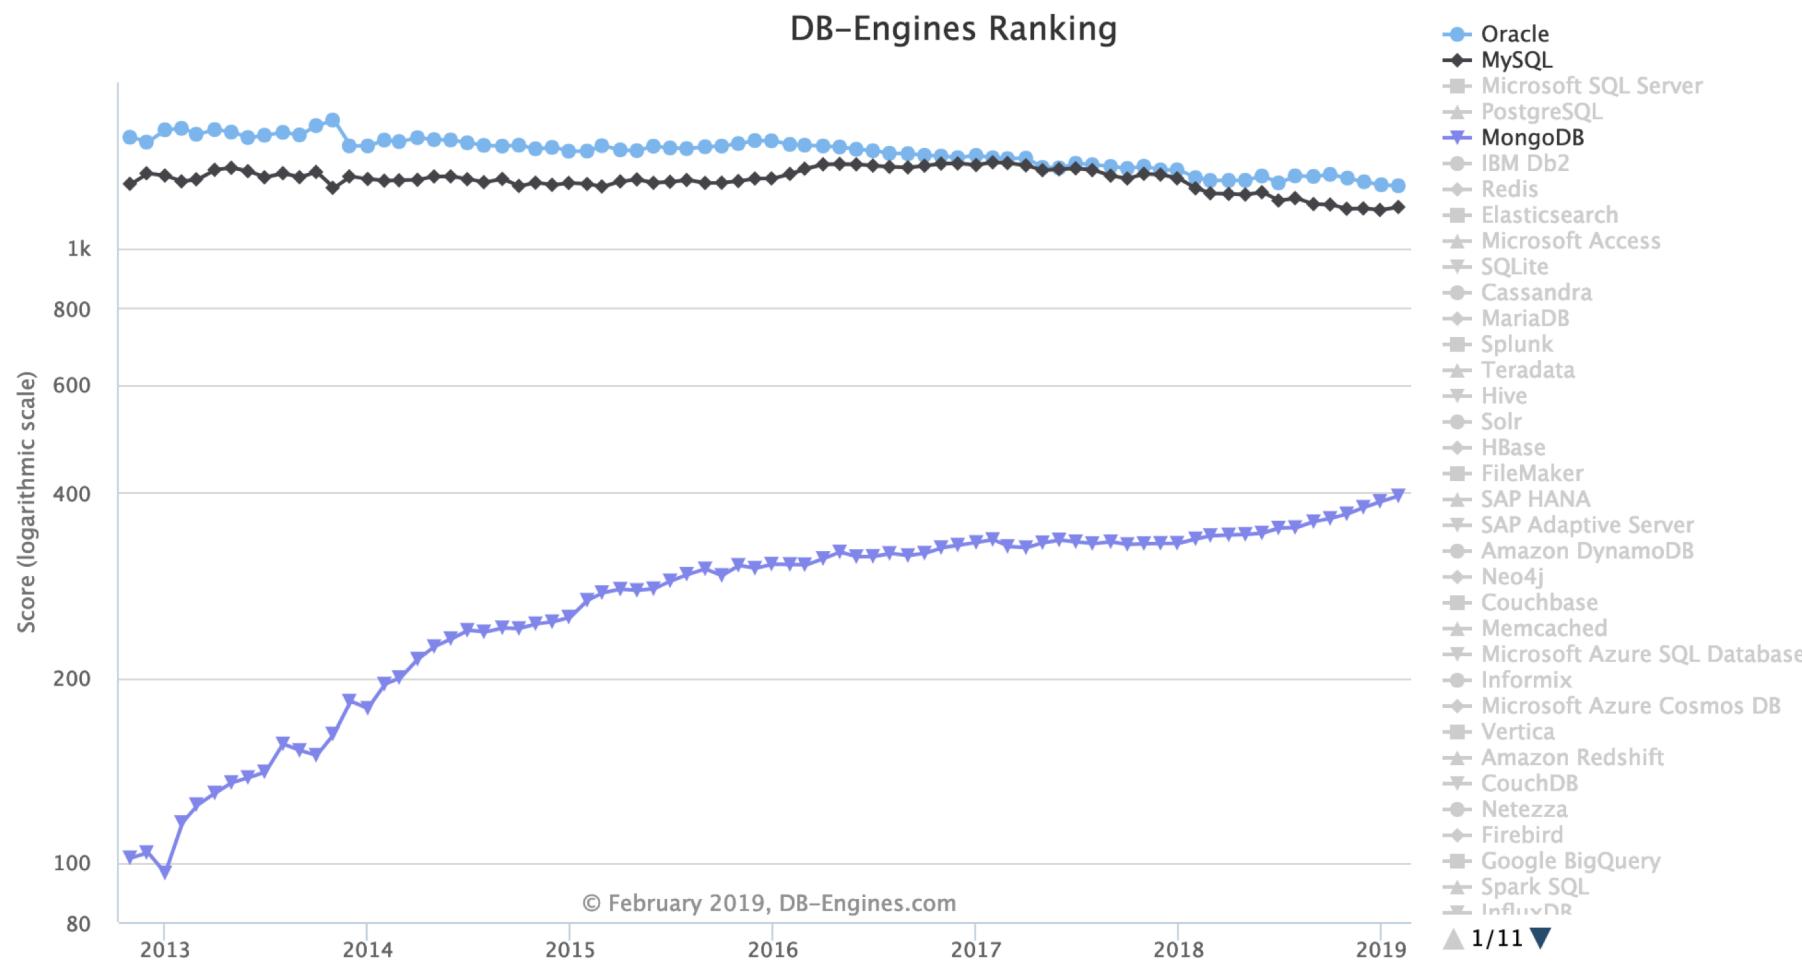
# Why are you here?

---

# DB-Engines Ranking - Trend Popularity

The DB-Engines Ranking ranks database management systems according to their popularity.

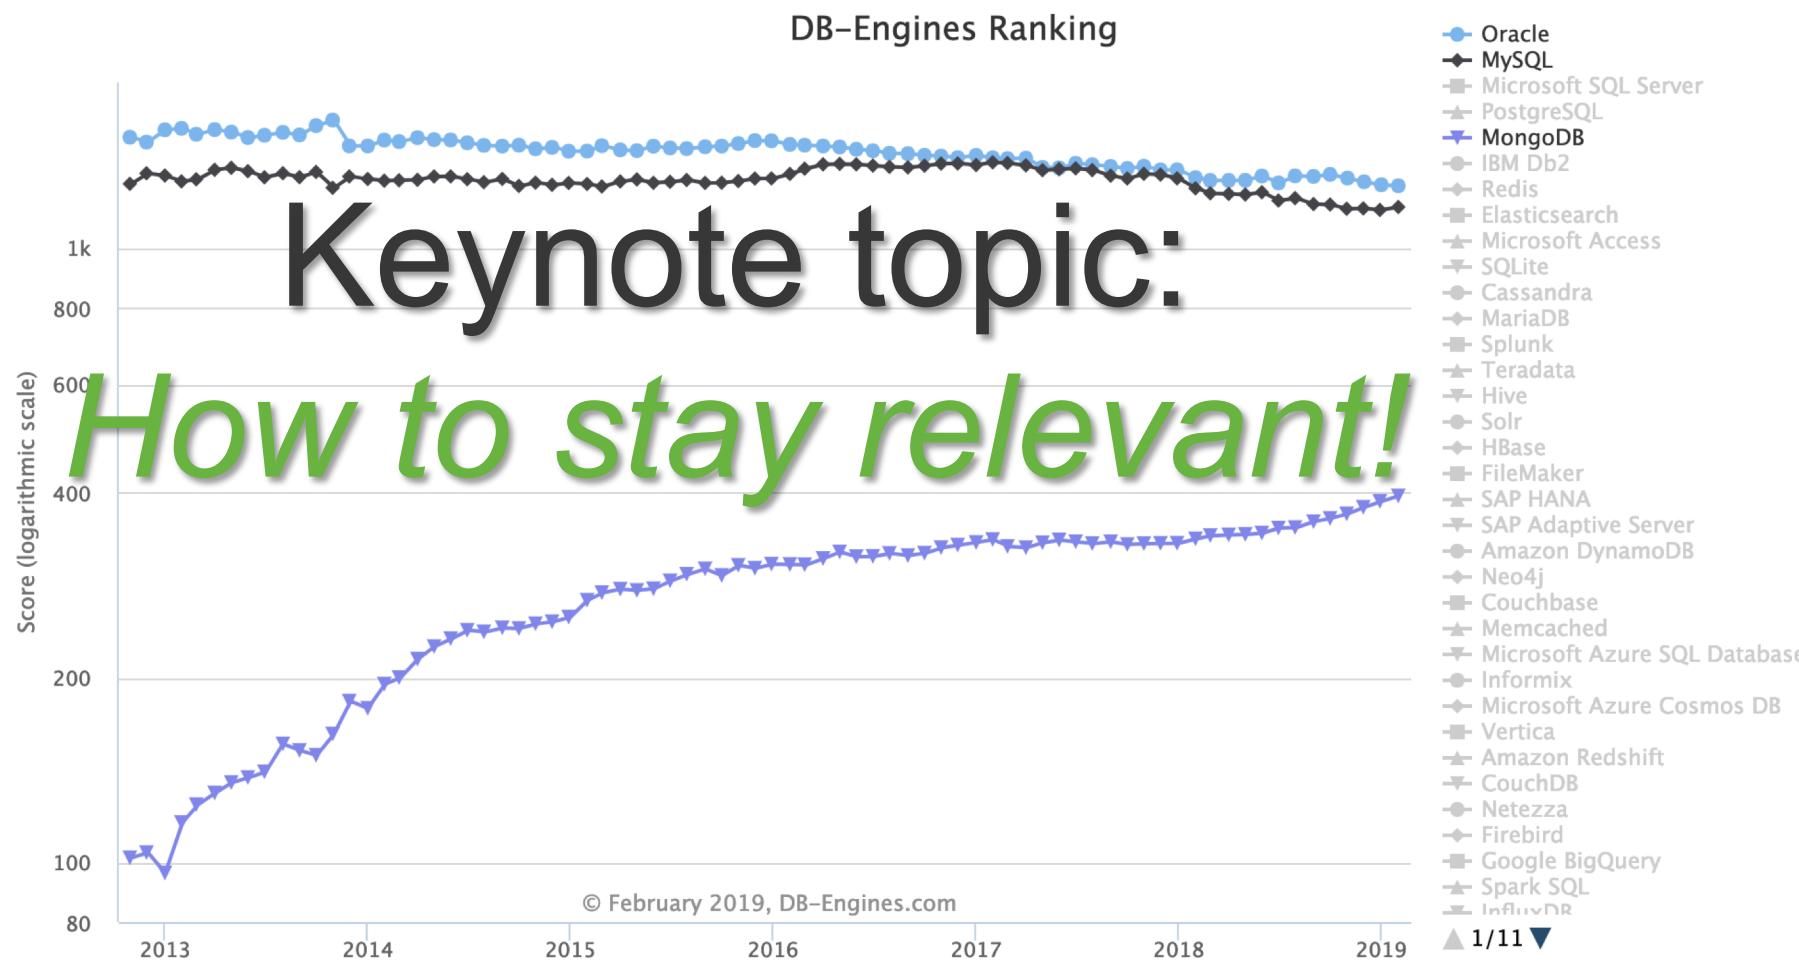
Read more about the [method](#) of calculating the scores.



# DB-Engines Ranking - Trend Popularity

The DB-Engines Ranking ranks database management systems according to their popularity.

Read more about the [method](#) of calculating the scores.





# MongoDB 101 Workshop

---

## Introduction to MongoDB (30 mins)

### Setting up MongoDB (30 mins)

- Laptop: Standalone Engine & Shell, Compass, MongoDB Driver (Python)
- Data Generator: MGenerateJS
- Cloud: MongoDB Atlas account

### Querying, Indexing and Aggregation (30 mins)

- Using Compass, MongoDB Shell & Python

### Schema Design & Use Cases (30 mins)

- Embedding & Referencing, 1-1, 1-M, M-M, Patterns
- Single View, IoT, Time Series, eCommerce/Content Management

### MongoDB Atlas (60 mins)

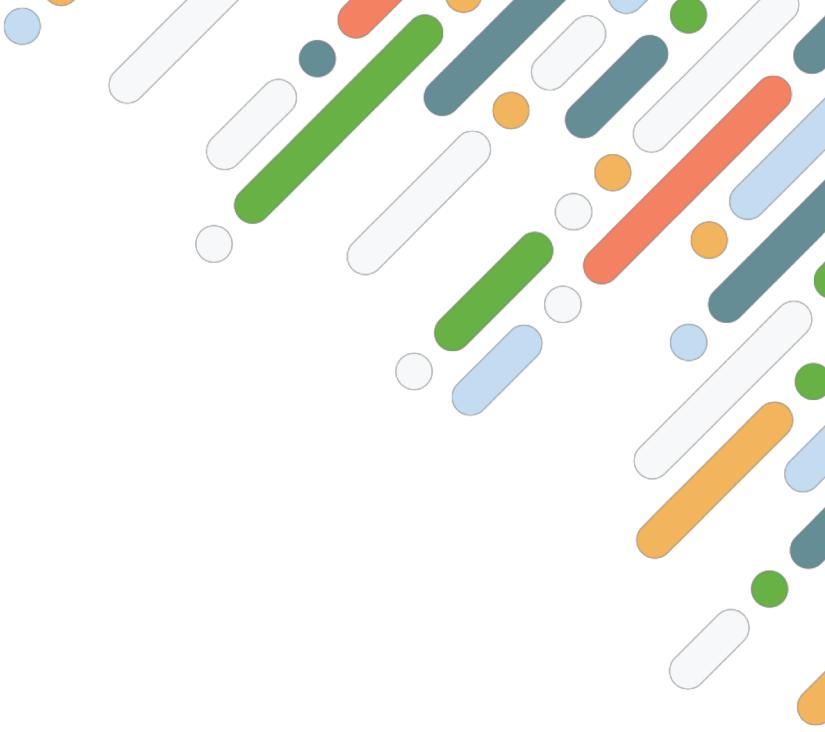
- Management, Clusters, Security
- High Availability & Scalability ( Replication & Sharding )
- Managed BI Connector & Charts
- ODBC for BI Connector
- MongoDB Stitch\*



# Introduction to MongoDB

9:00 – 9:30

---



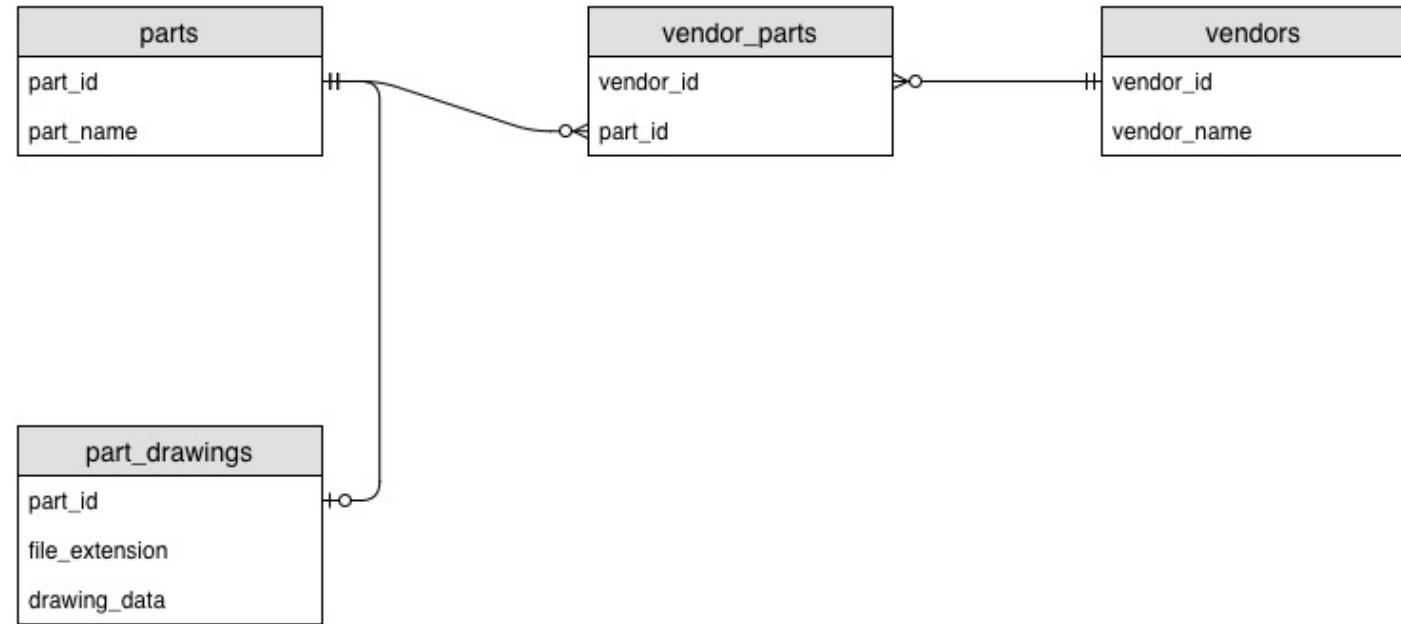
# Introduction to MongoDB

---



# Relational Database Design

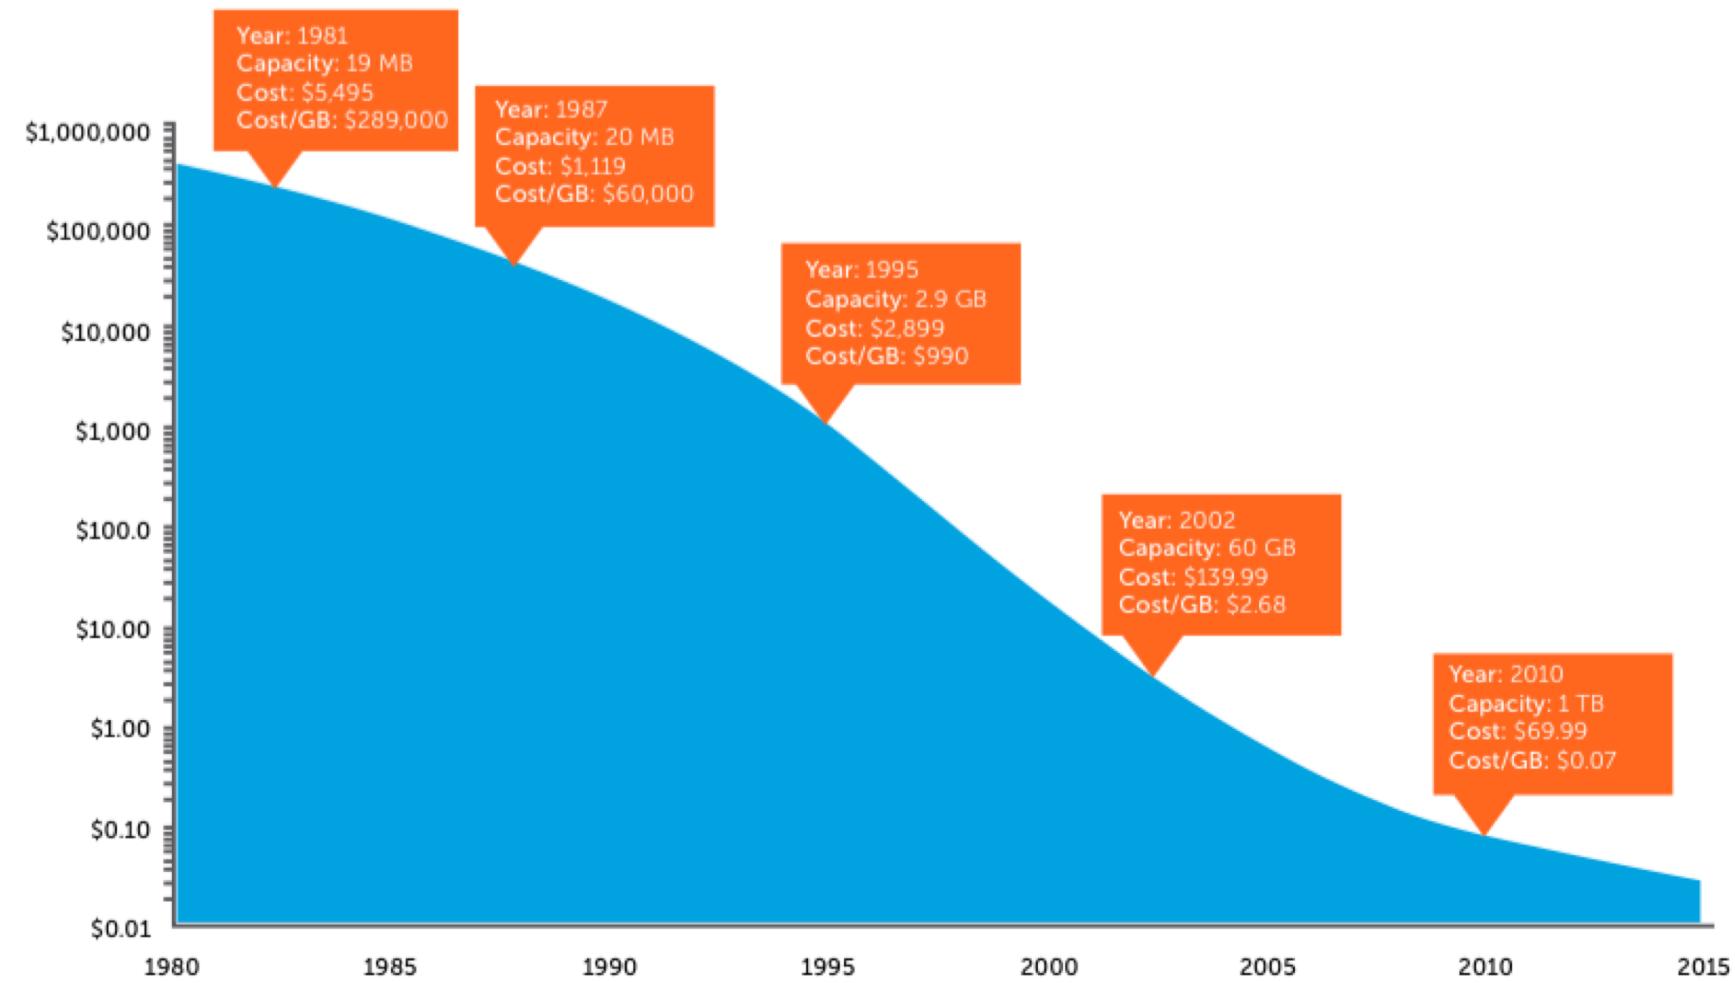
---



# Why Did We Normalize Data?

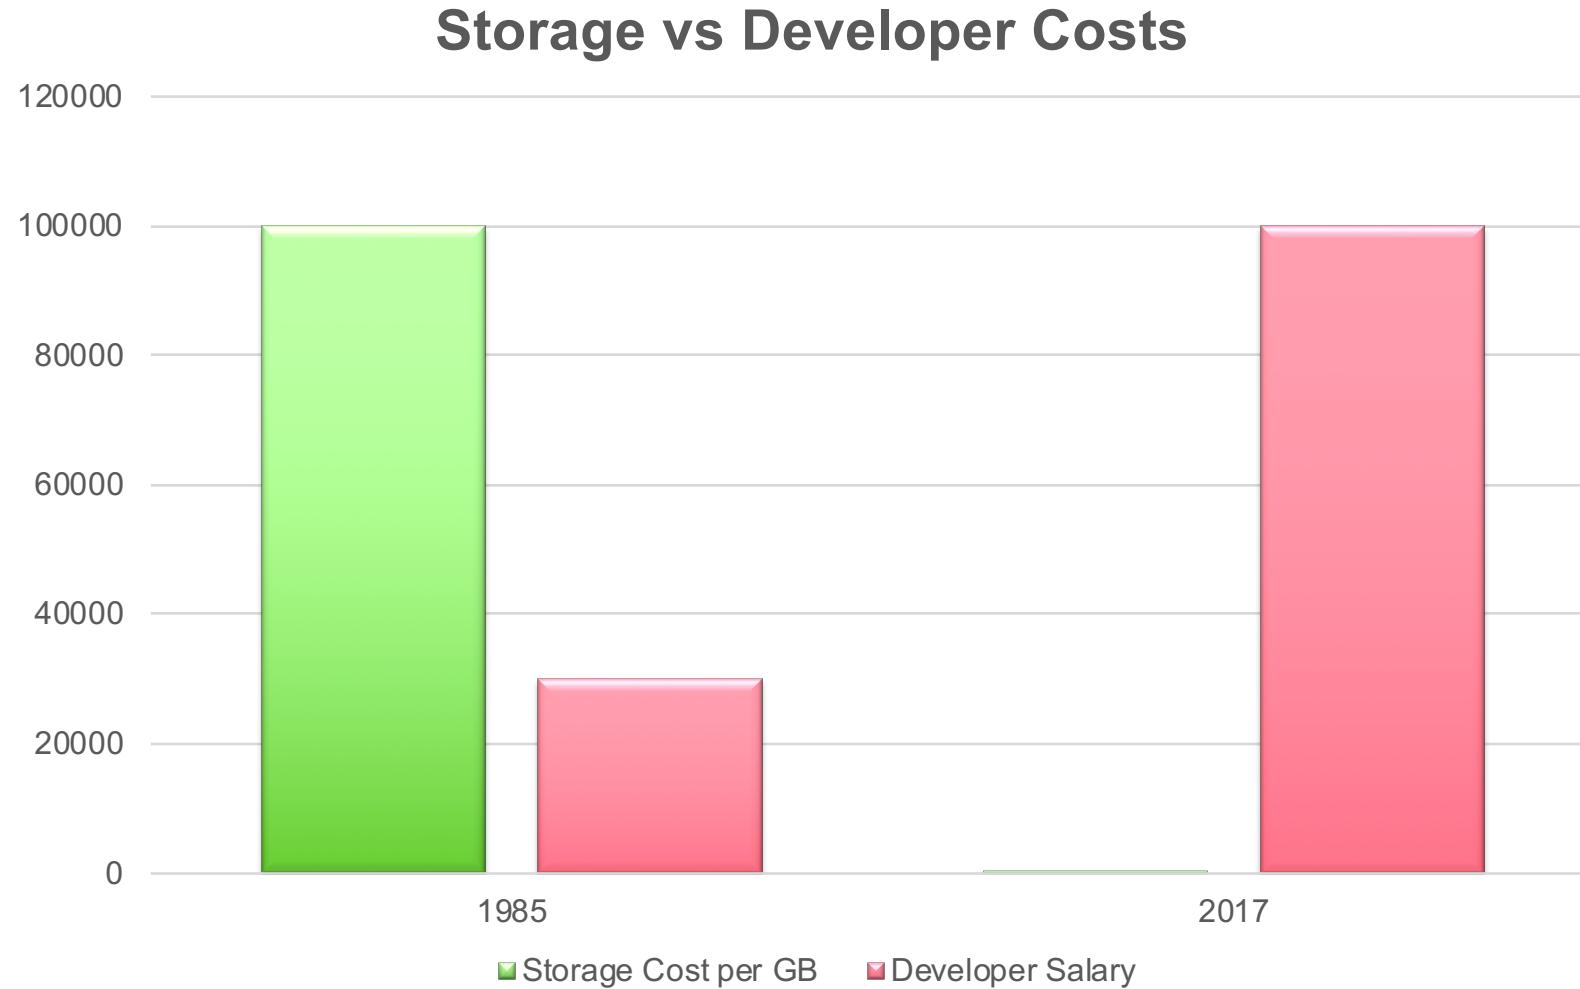
---

Storage cost  
over time



# Costs have Shifted! – Developers are the new kingmakers!

---





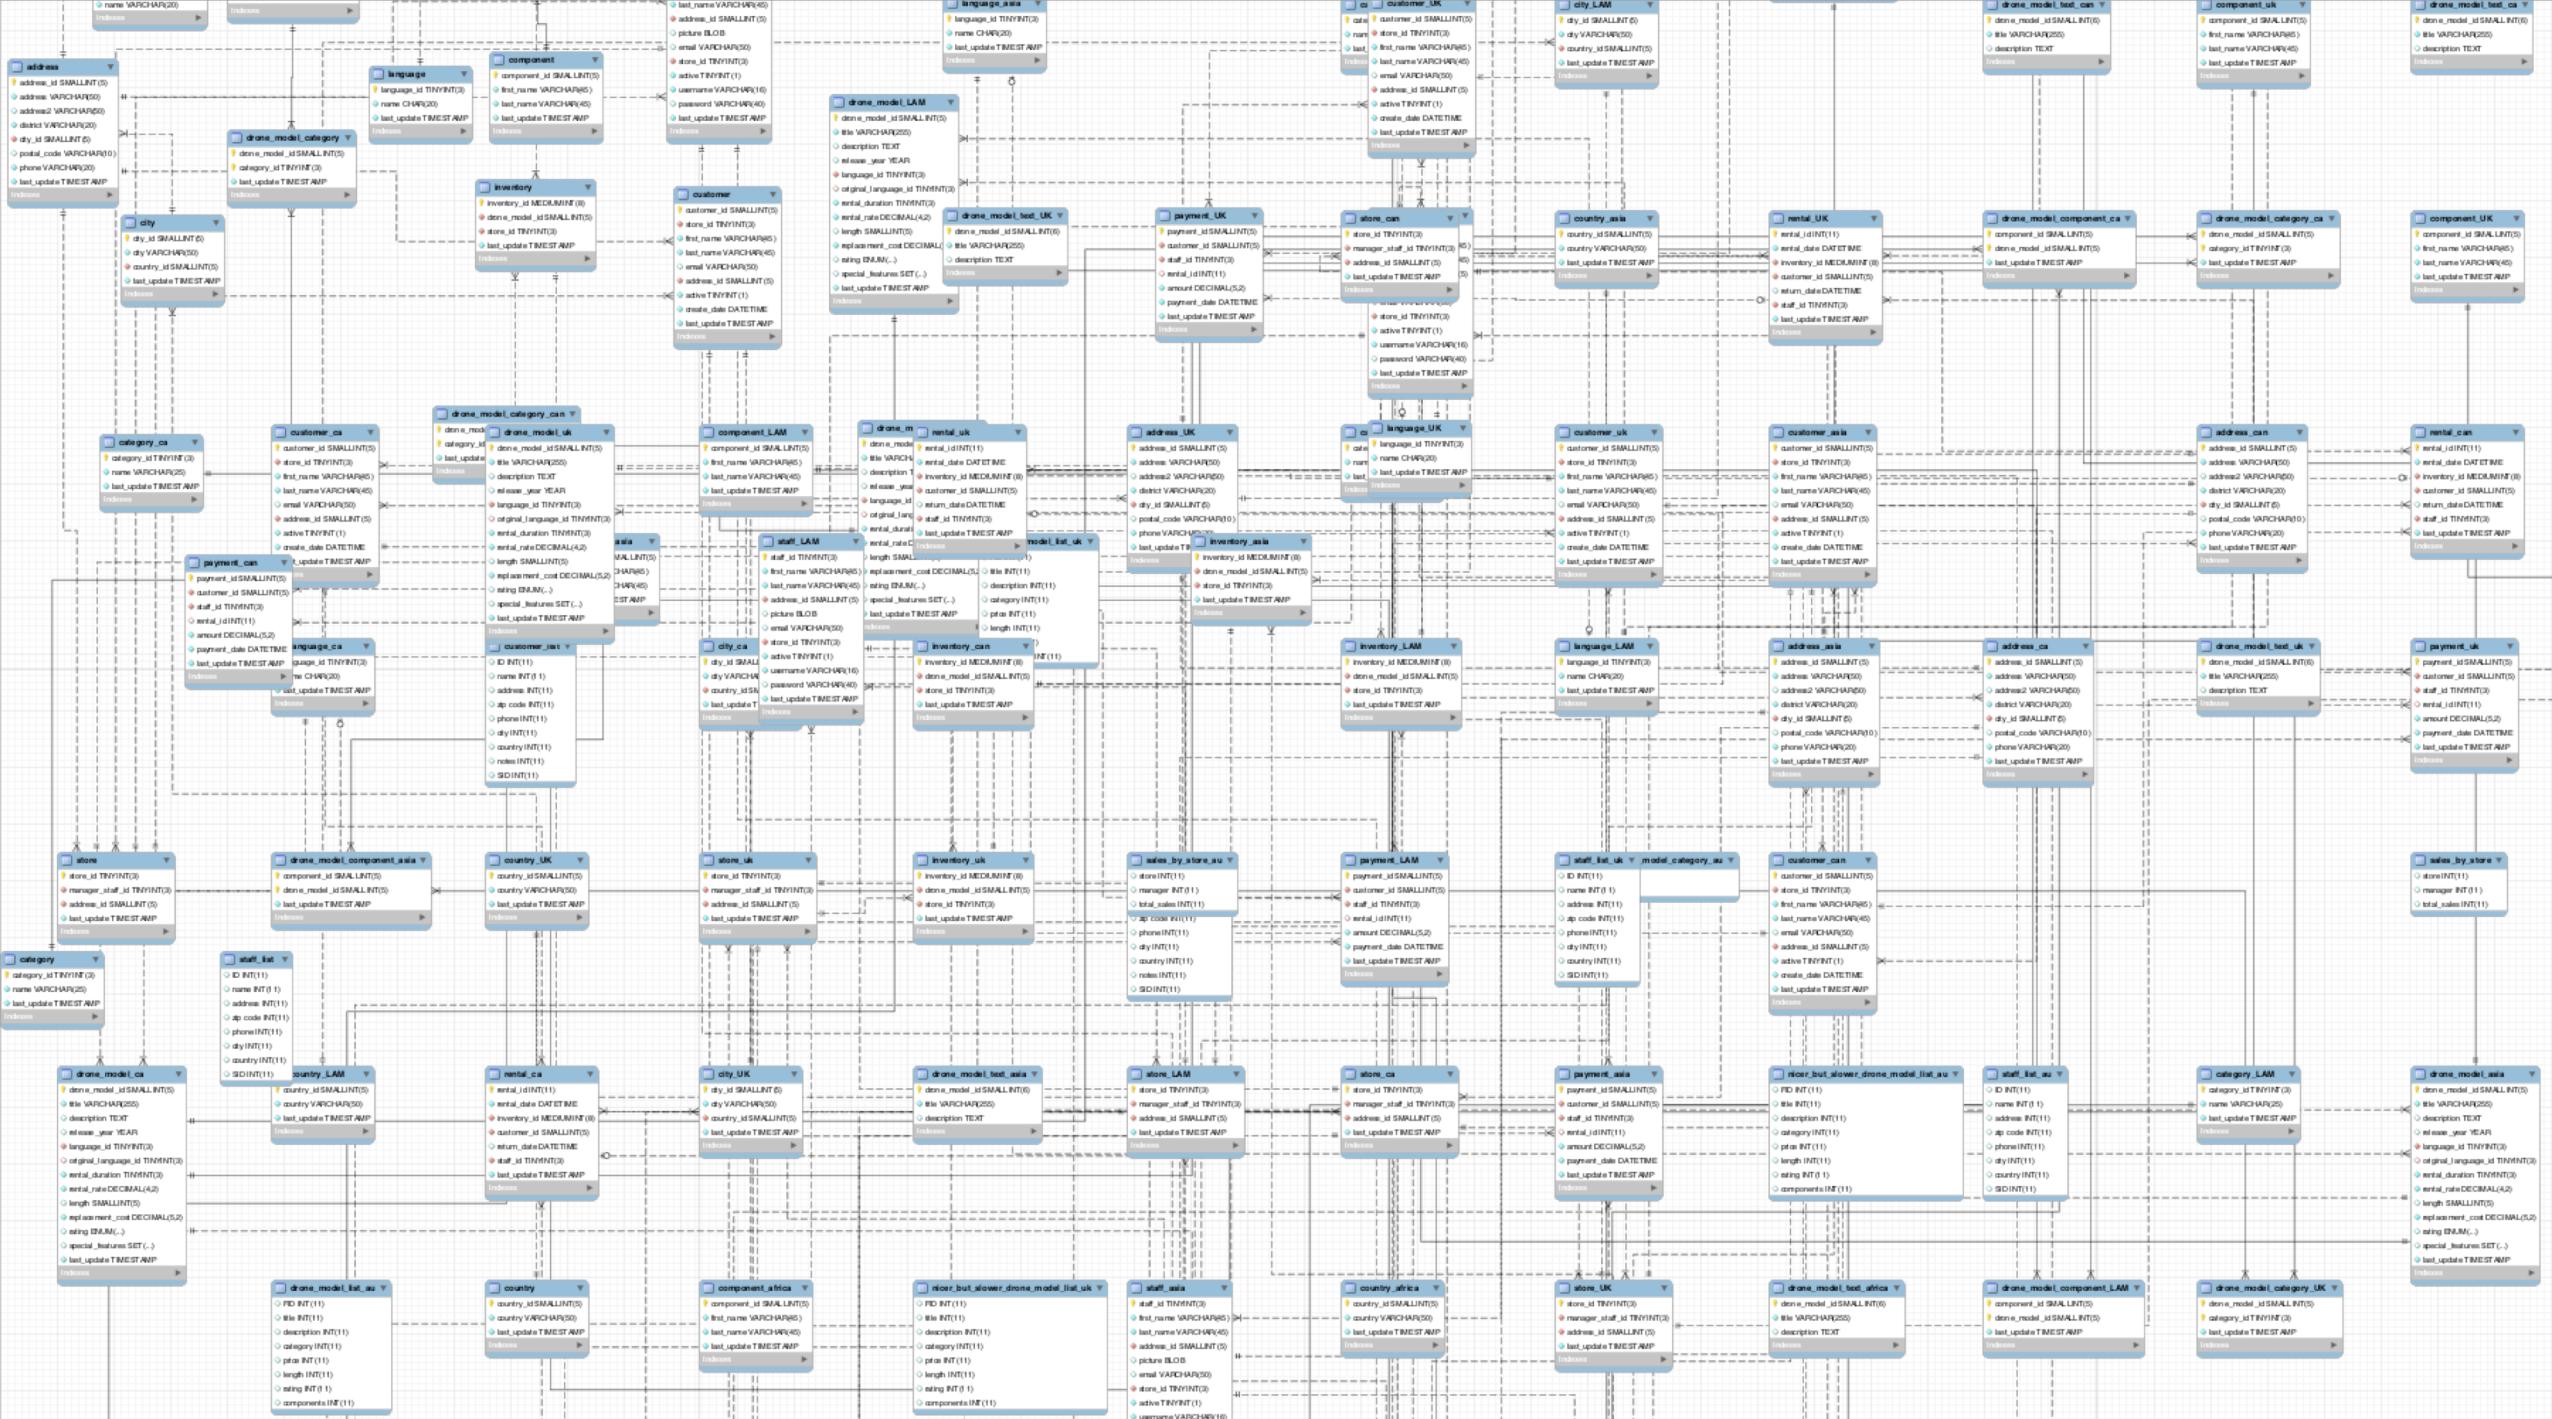
# Tabular/Relational

---

vs.

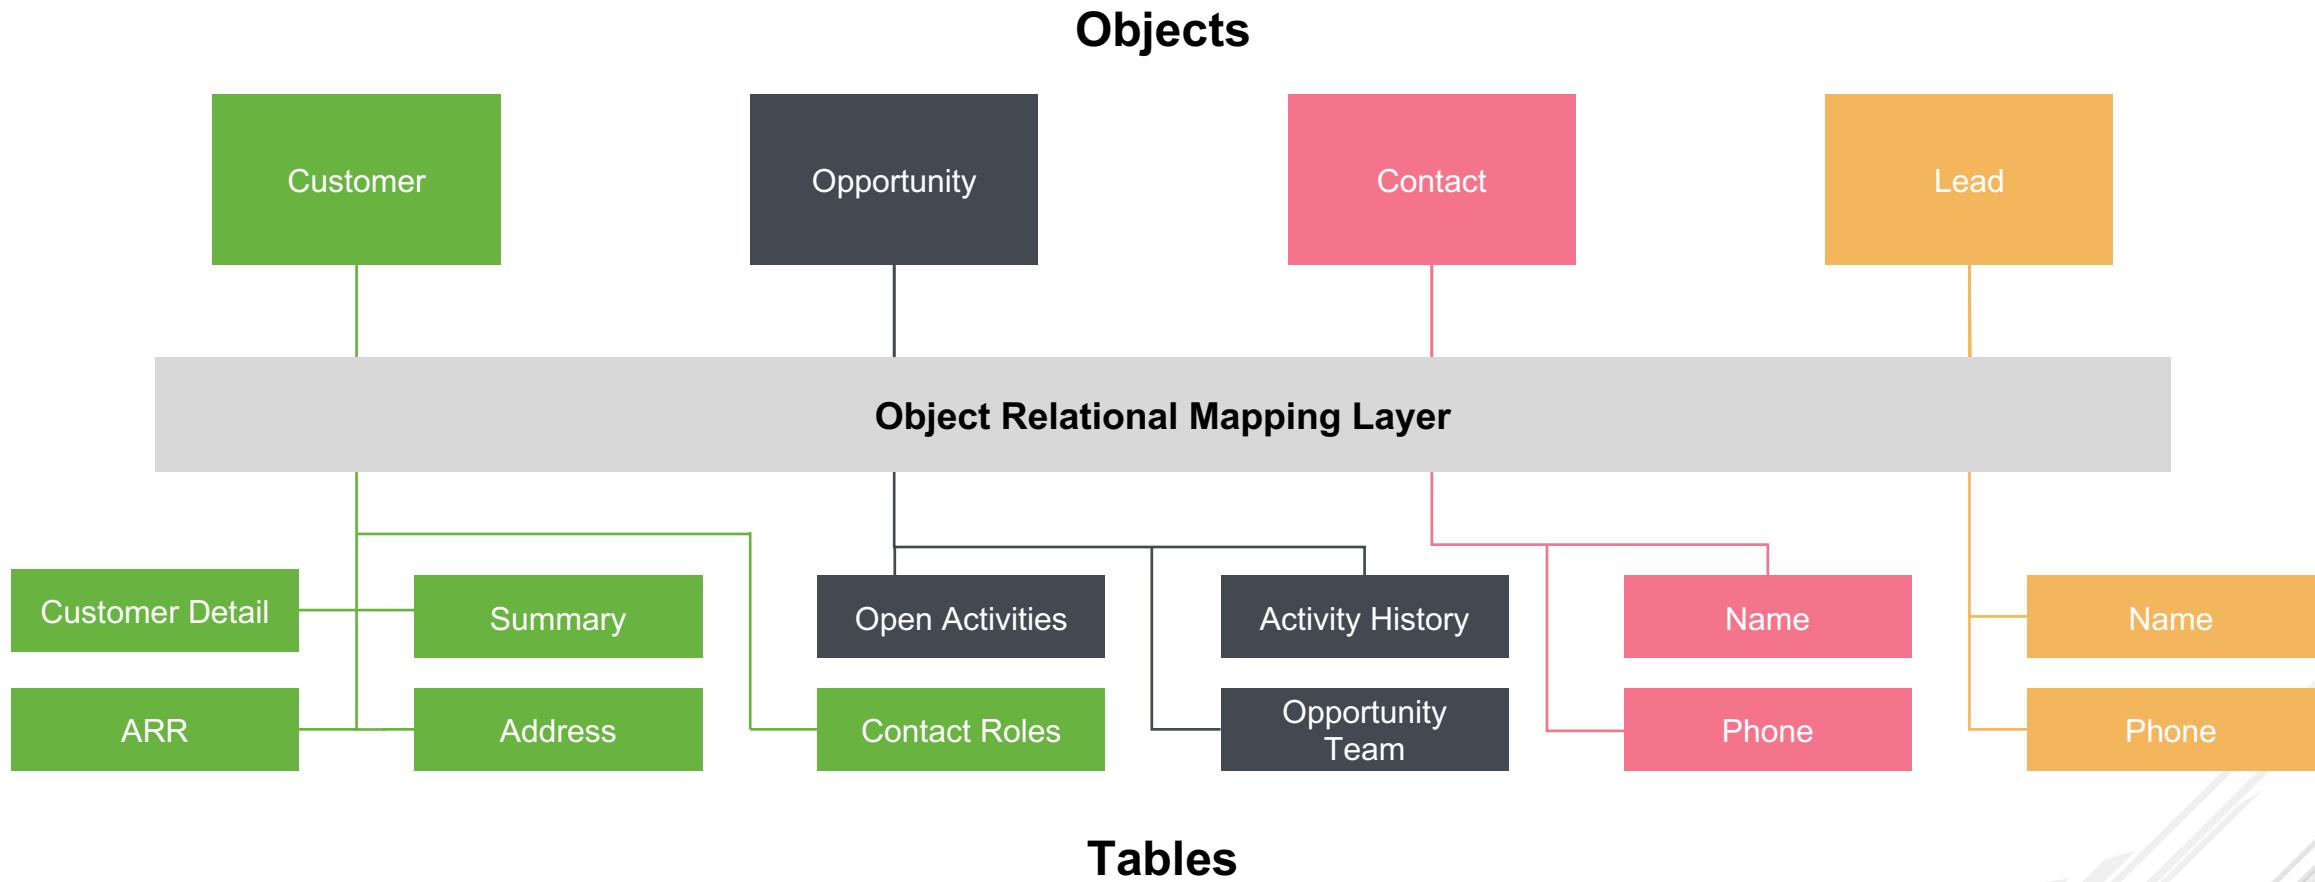
# Object Oriented





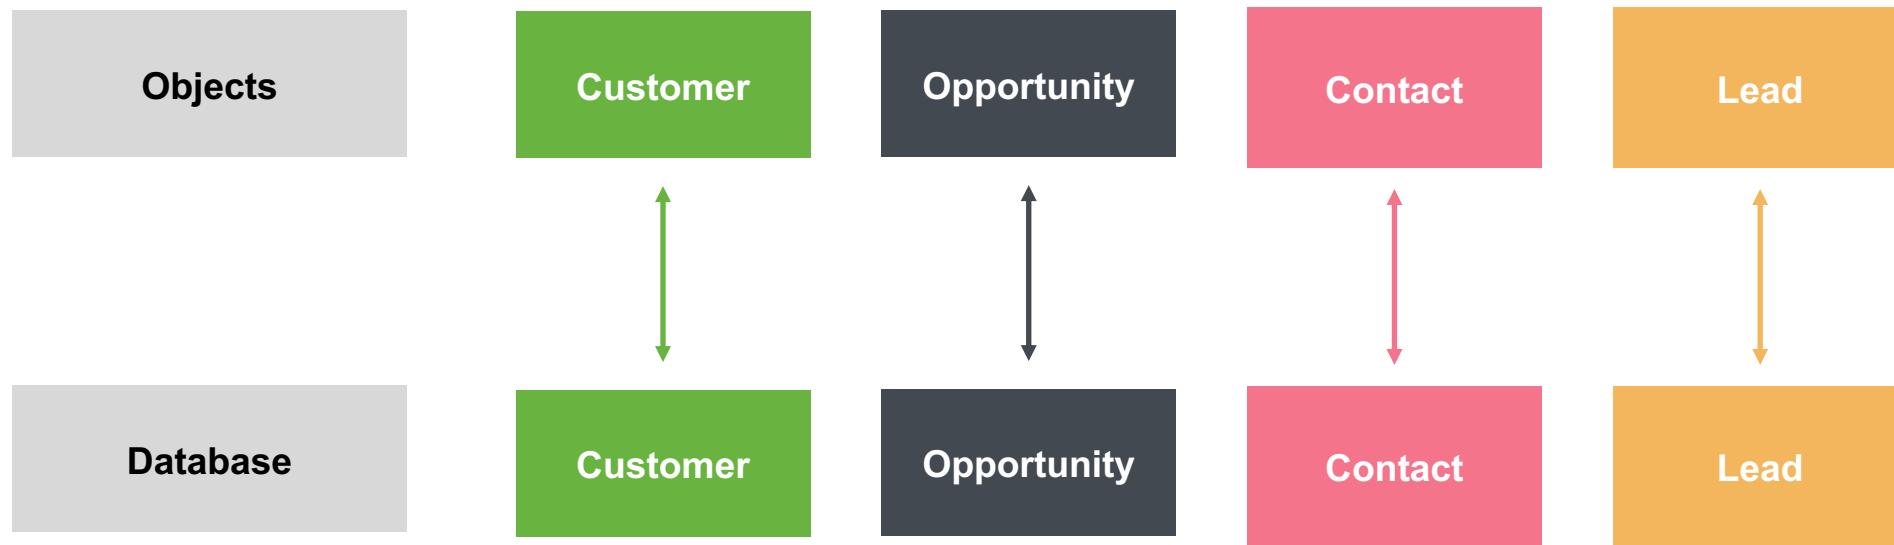


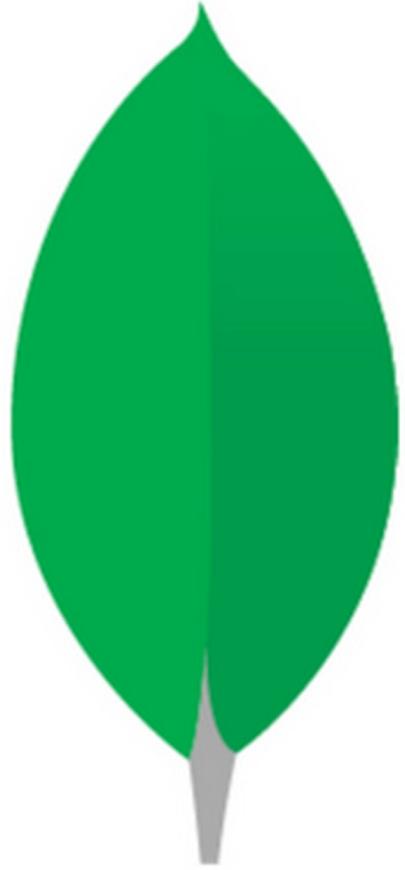
# Go from this....





To this: store objects directly...





Our founders believe that coding should be natural, and so should using a database.

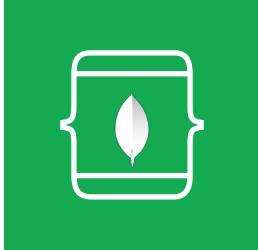
They want the experience of using MongoDB to be simple and natural. Thus, the leaf.



# Why MongoDB?

---

## Intelligent Operational Data Platform



**Best way to work  
with data**



**Intelligently put data  
where you need it**



**Freedom  
to run anywhere**

# The evolution of MongoDB

	3.0	3.2	3.4	3.6	4.0
Doc-Level Concurrency					Multi-Document ACID Transactions
Compression					Atlas Global Clusters
Storage Engine API					Atlas HIPAA
≤50 replicas					Atlas LDAP
Auditing ++					Atlas Audit
Ops Manager					Atlas Encrypted Storage Engine
	Document Validation	Linearizable reads	Change Streams		Atlas AWS Backup Snapshots
	\$lookup	Intra-cluster compression	Retryable Writes		Atlas Full CRUD
	Fast Failover	Views	Expressive Array Updates		Agg Pipeline Type Conversions
	Simpler Scalability	Log Redaction	Query Expressivity		40% Faster Shard Migrations
	Aggregation ++	Graph Processing	Causal Consistency		Snapshot Reads
	Encryption At Rest	Decimal	Consistent Sharded Sec. Reads		Non-Blocking Secondary Reads
	In-Memory Storage Engine	Collations	Compass Community		SHA-2
	BI Connector	Faceted Navigation	Ops Manager ++		TLS 1.1+
	MongoDB Compass	Zones ++	Query Advisor		Compass Agg Pipeline Builder
	APM Integration	Aggregation ++	Schema Validation		Compass Export to Code
	Profiler Visualization	Auto-balancing ++	End to End Compression		Charts Beta
	Auto Index Builds	ARM, Power, zSeries	IP Whitelisting		Free Monitoring Cloud Service
	Backups to File System	BI & Spark Connectors ++	Default Bind to Localhost		Ops Manager K8s & OpenShift
		Compass ++	Sessions		MongoDB Stitch GA
		Hardware Monitoring	WiredTiger 1m+ Collections		MongoDB Mobile GA
		Server Pool	MongoDB BI Connector ++		
		LDAP Authorization	Expressive \$lookUp		
		Encrypted Backups	R Driver		
		Cloud Foundry Integration	Atlas Cross Region Replication		
			Atlas Auto Storage Scaling		



# MongoDB Use Cases

## Single View

- Real-time views of your business that integrate all of your siloed data.
- Flexible Data Model

## Internet of Things

- Variety of data – multiple versions and schemas over time – older & newer models
- Horizontal Scale

## Mobile

- JSON Payloads – single fetch
- Sync from Data Center to Device

## Real Time Analytics

- Analyze on the spot – no ETL / DW
- MQL, SQL or ML/Data Science

## Personalization

## Catalog

## Content Management

<https://www.mongodb.com/use-cases>

# Best Way to Work With Data

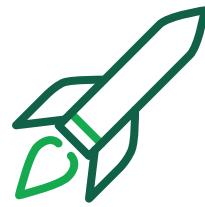
---



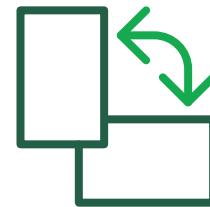
**Easy:**  
Work with data  
in a natural,  
intuitive way



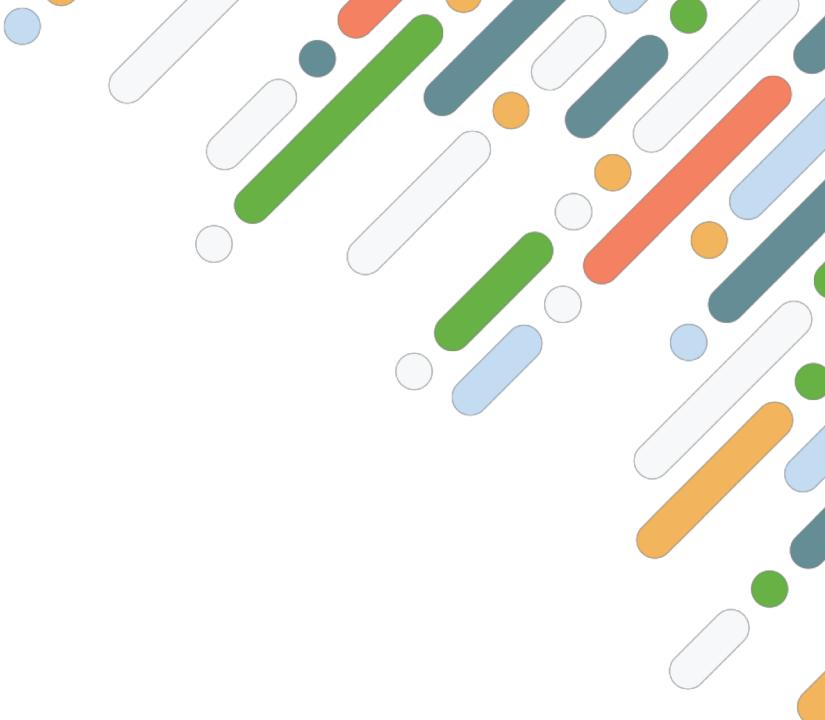
**Flexible:**  
Adapt and  
make changes  
quickly



**Fast:**  
Get great  
performance  
with less code



**Versatile:**  
Supports a  
wide variety of  
data models  
and queries



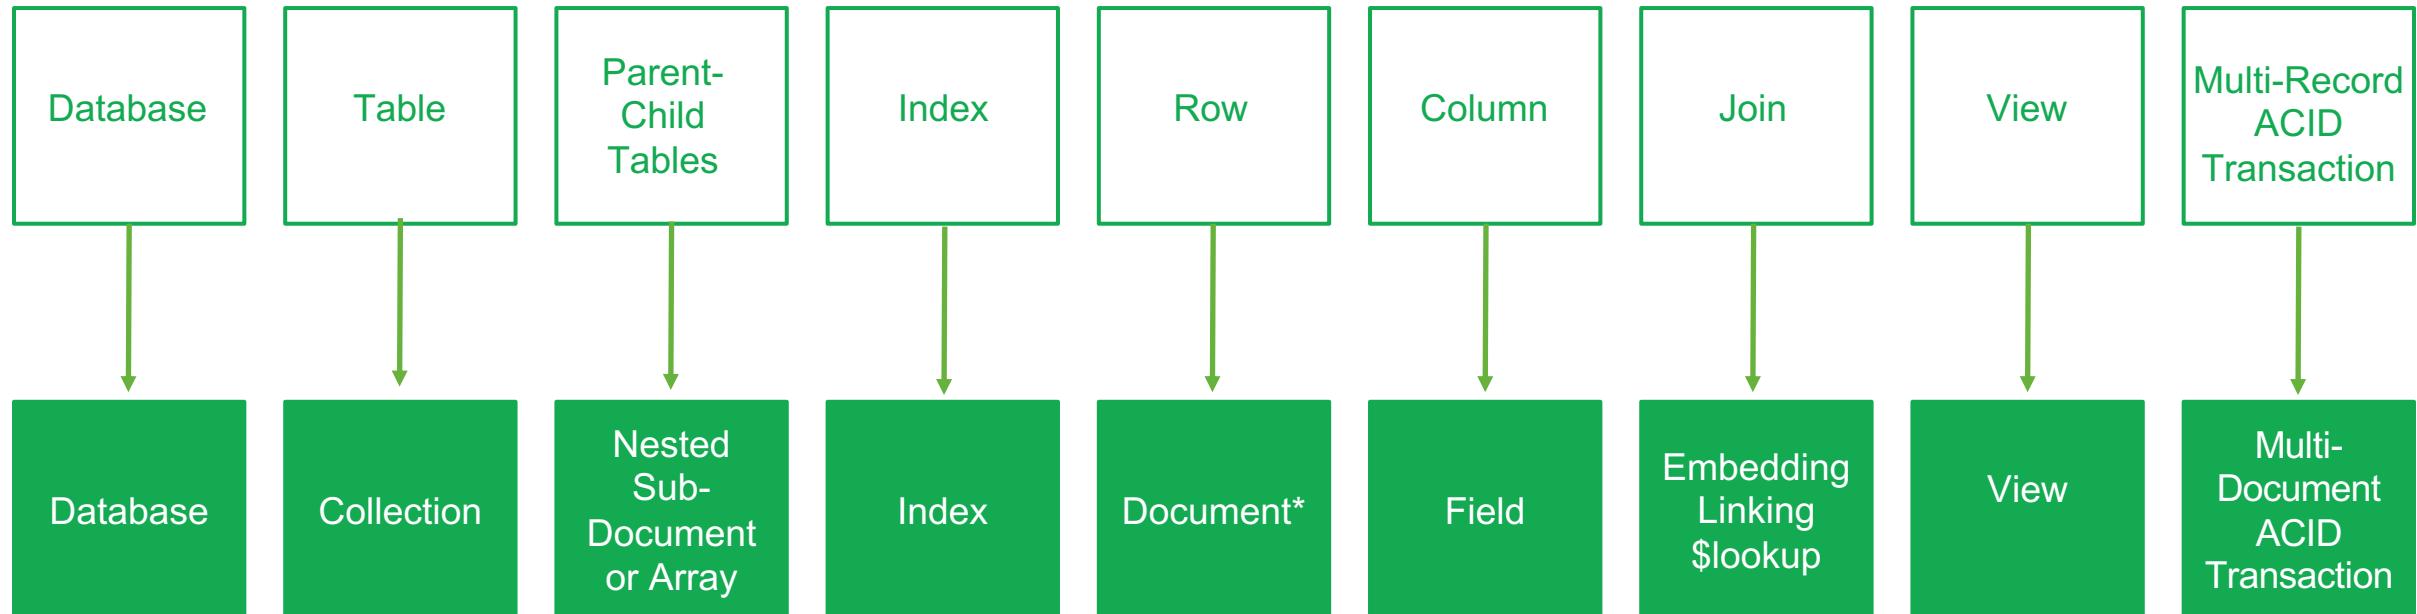
# The Document Model

---

# Terminology

---

## RDBMS



## MongoDB 4.0

\* Proper document schema design yields more entity data per document than found in a relational database row



## JSON is Widely Used

---

JSON is fast, lean, flexible

JSON is supported by almost all languages and tools

JSON is used for modern communication and data exchange

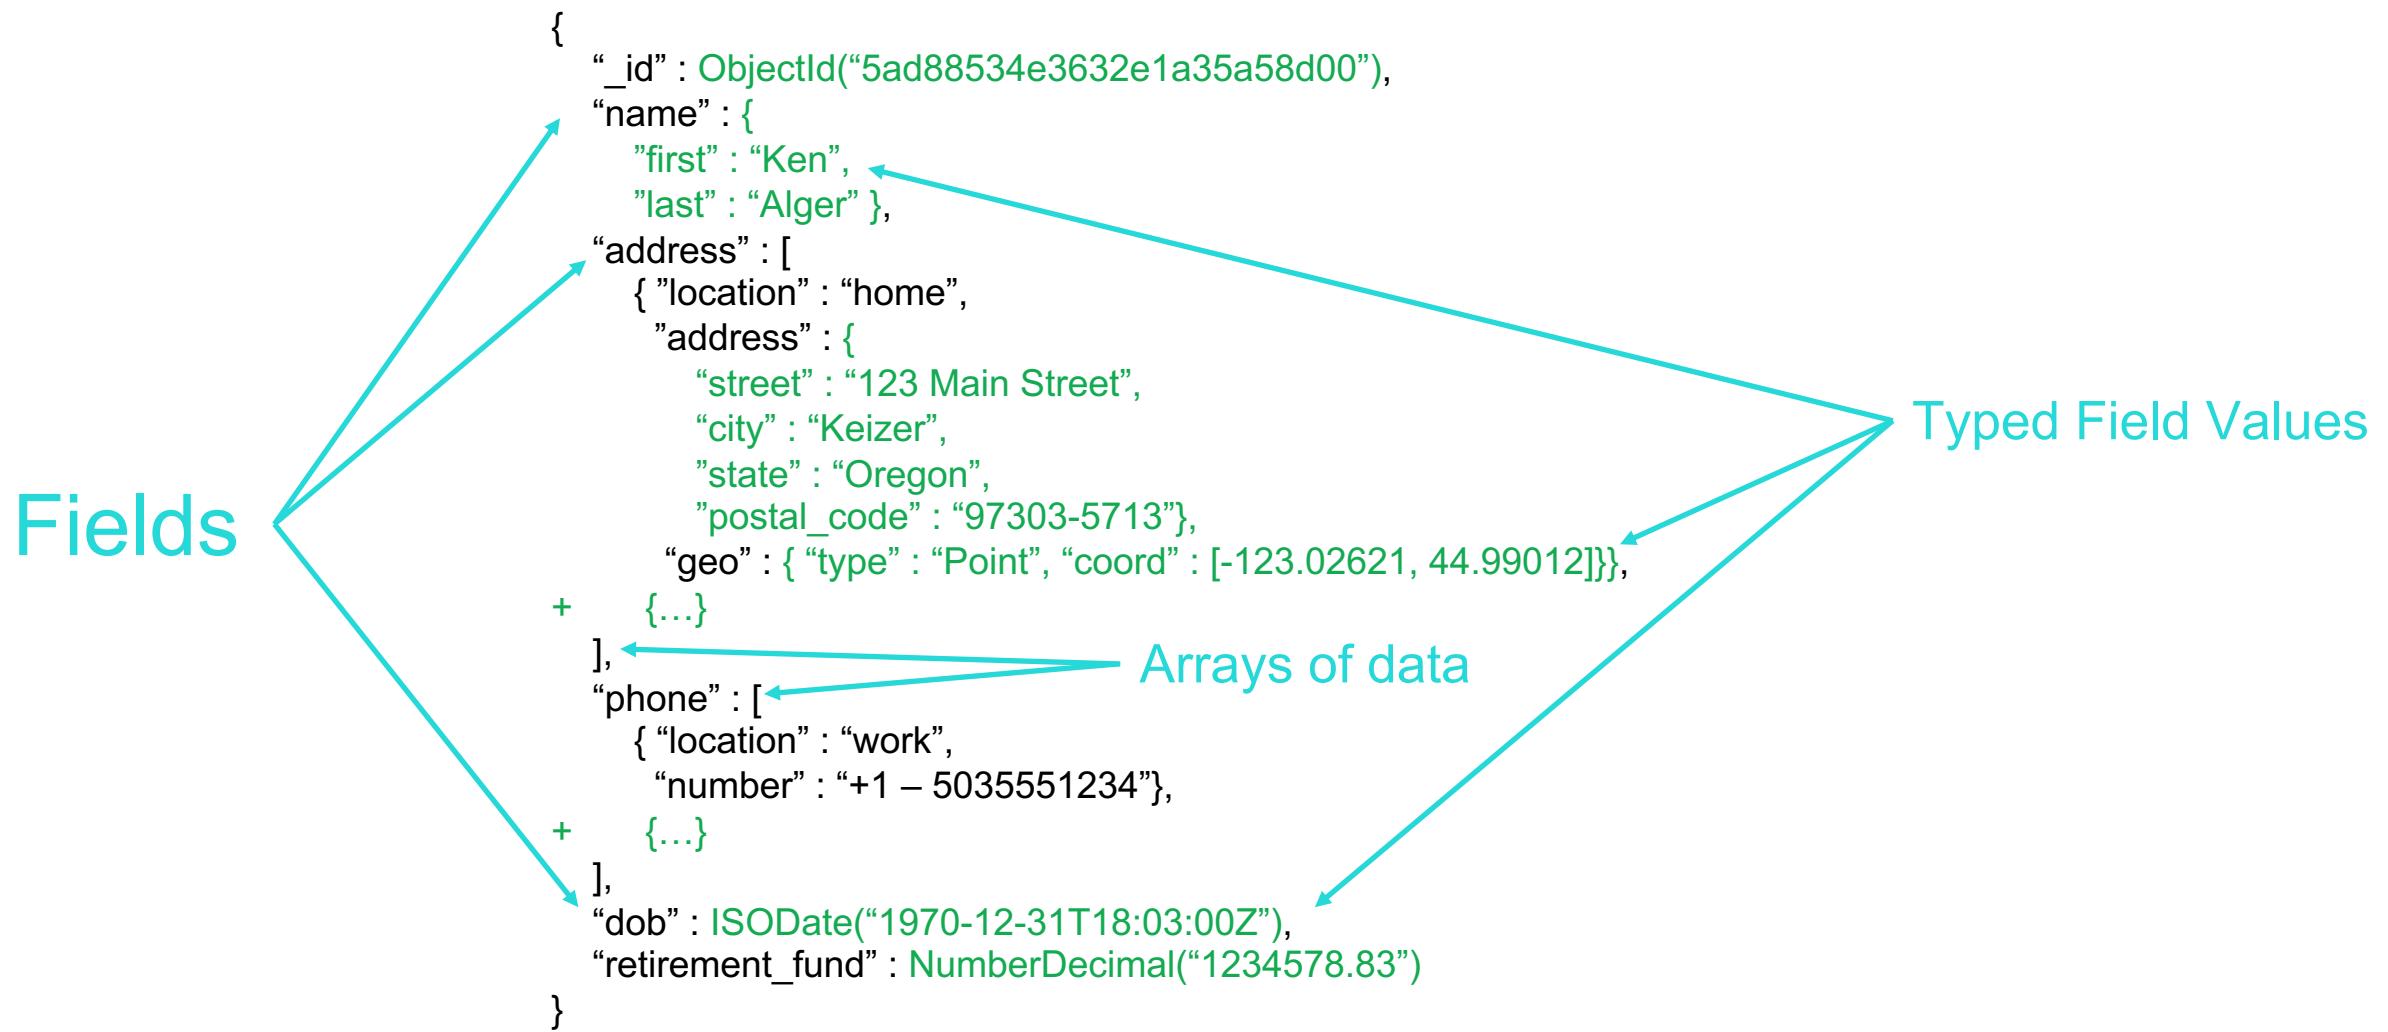
# Rich Data Similar to JSON

---



Binary serialization of JSON data representation.  
Easier to transfer than text files.  
Lightweight, transversible, efficient.

# Document Data Model





## Document Data Model

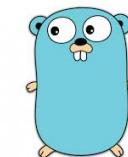
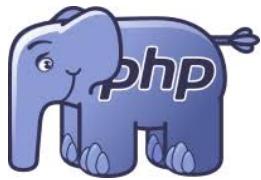
---

- Naturally maps objects to code – **developers think in terms of objects**, why shouldn't your data be modeled that way?
- Represent data of any structure
- Strongly typed for ease of processing – Over 20 binary encoded **JSON data types**
- Access by idiomatic **drivers in all major programming languages**

# Drivers & Frameworks

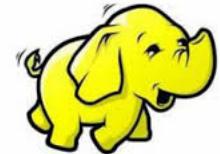
---

## Drivers



(Beta)

## Frameworks





# I've heard MongoDB is NoSQL – Does that mean I can't query?

---

MongoDB Query Language

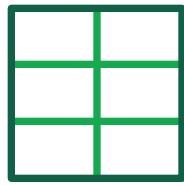


# Multiple data models, rich query functionality

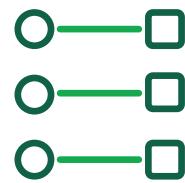
---



*JSON Documents*



*Tabular*



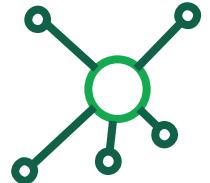
*Key-Value*



*Text*



*Geospatial*



*Graph*

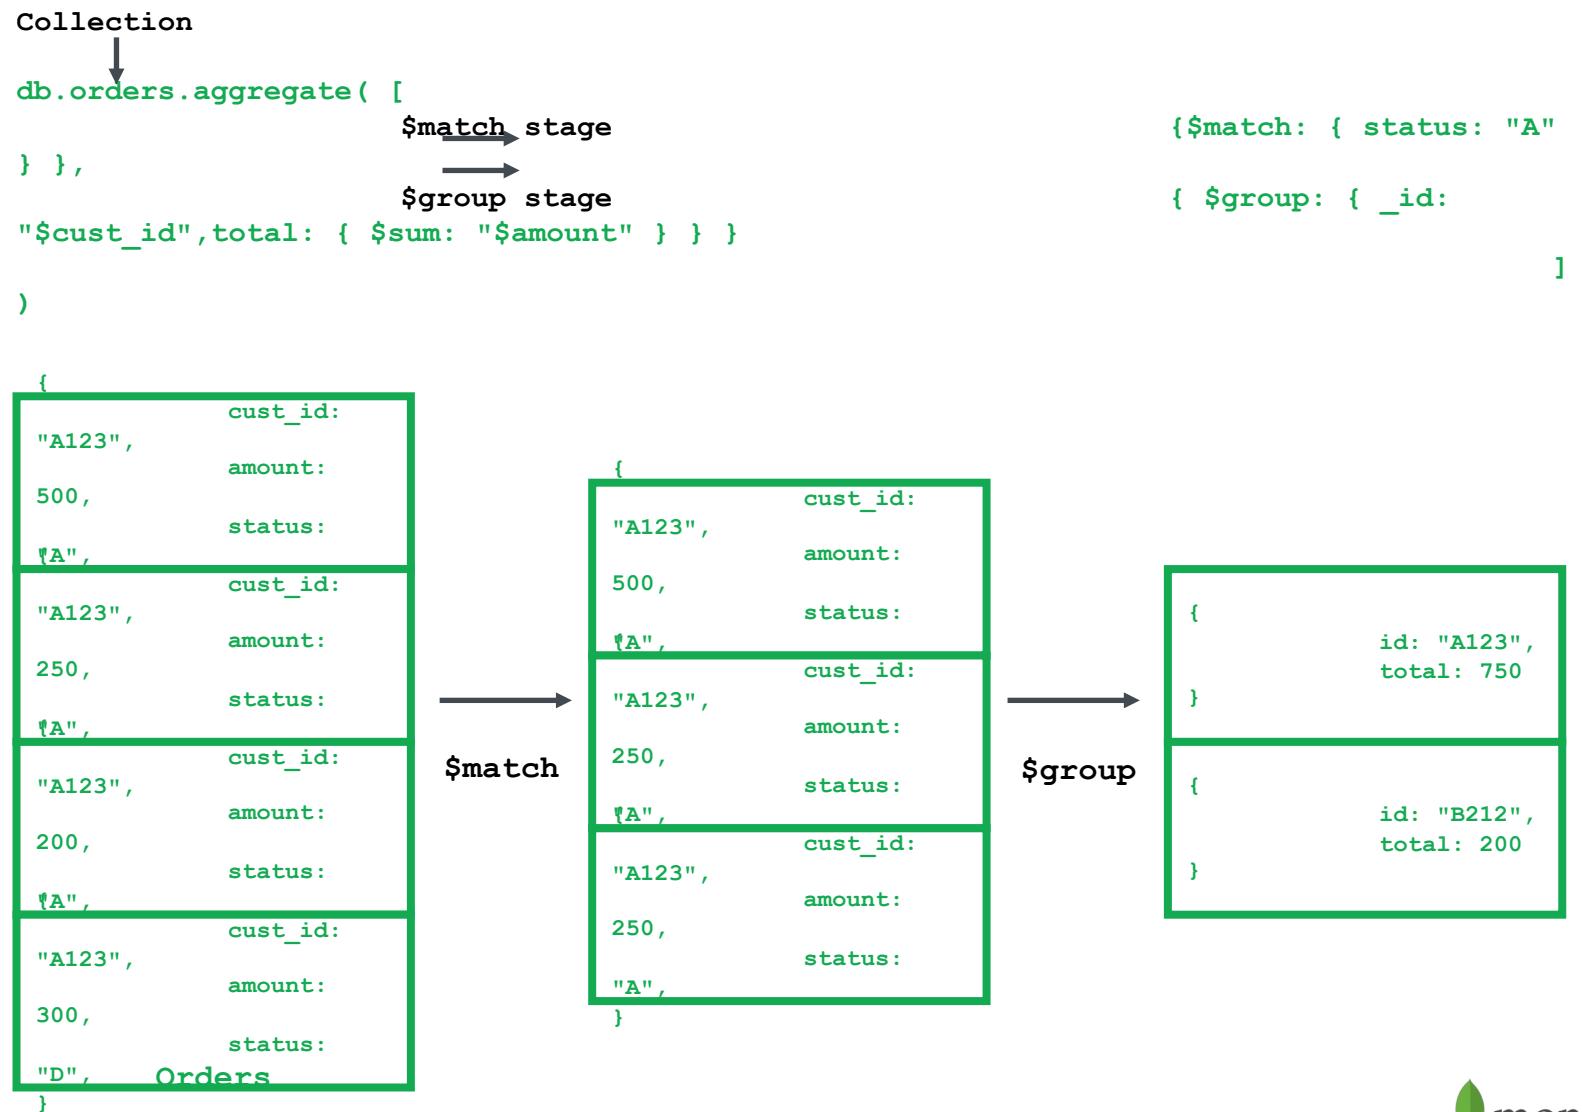
## Rich Queries

Point | Range | Geospatial | Faceted Search | Aggregations | JOINS | Graph Traversals

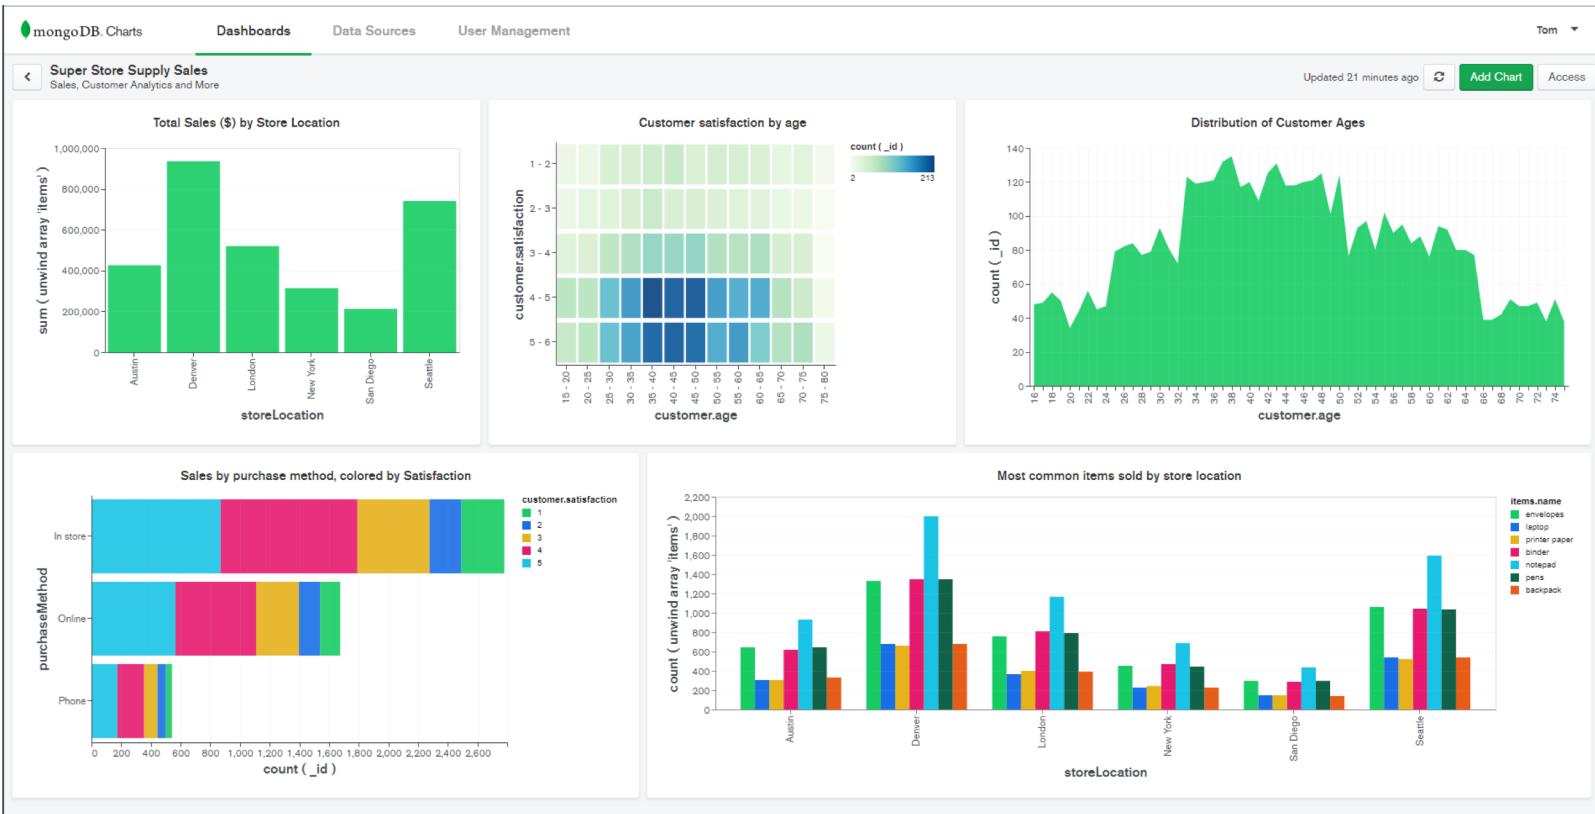
# Aggregation Framework

Advanced data processing pipeline for transformations and analytics

- Multiple stages
- Similar to a unix pipe
  - Build complex pipeline by chaining commands together
- Rich Expressions



# Sophisticated Analytics & Visualizations Of Data In Place

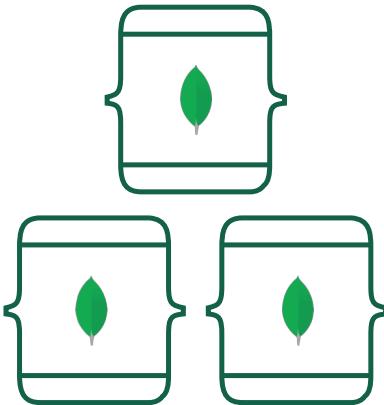


- Rich MongoDB query language & idiomatic drivers
- Connector for BI
- Connector for Spark
- Charts (beta)



# Data Availability

---

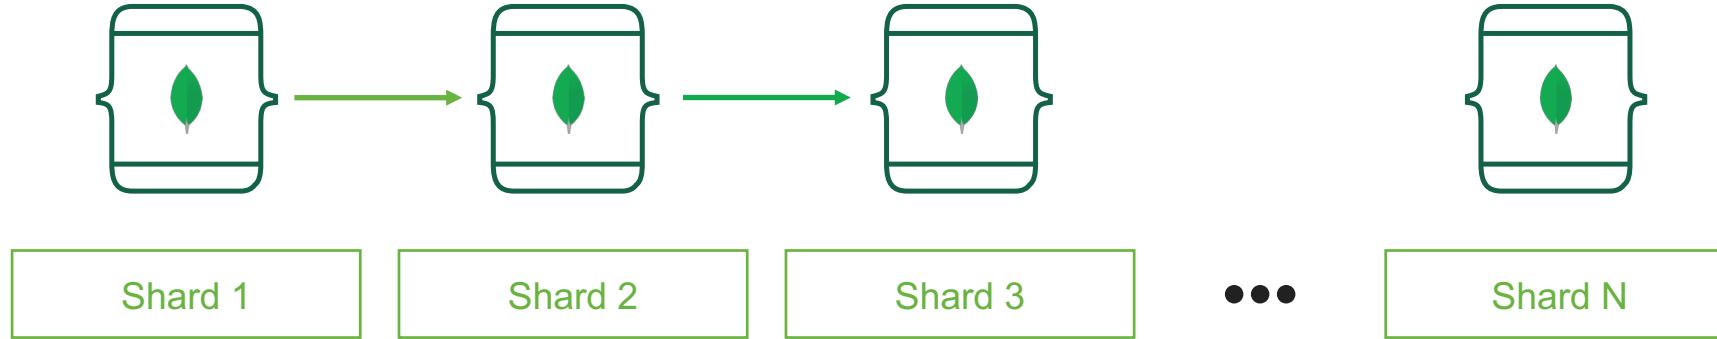


## Replica Sets

- Up to 50 replicas, distributed across racks, data centers, and regions
- Self-healing
- Tunable durability and consistency controls
- Always-on write availability with retryable writes

# Data Scalability

---

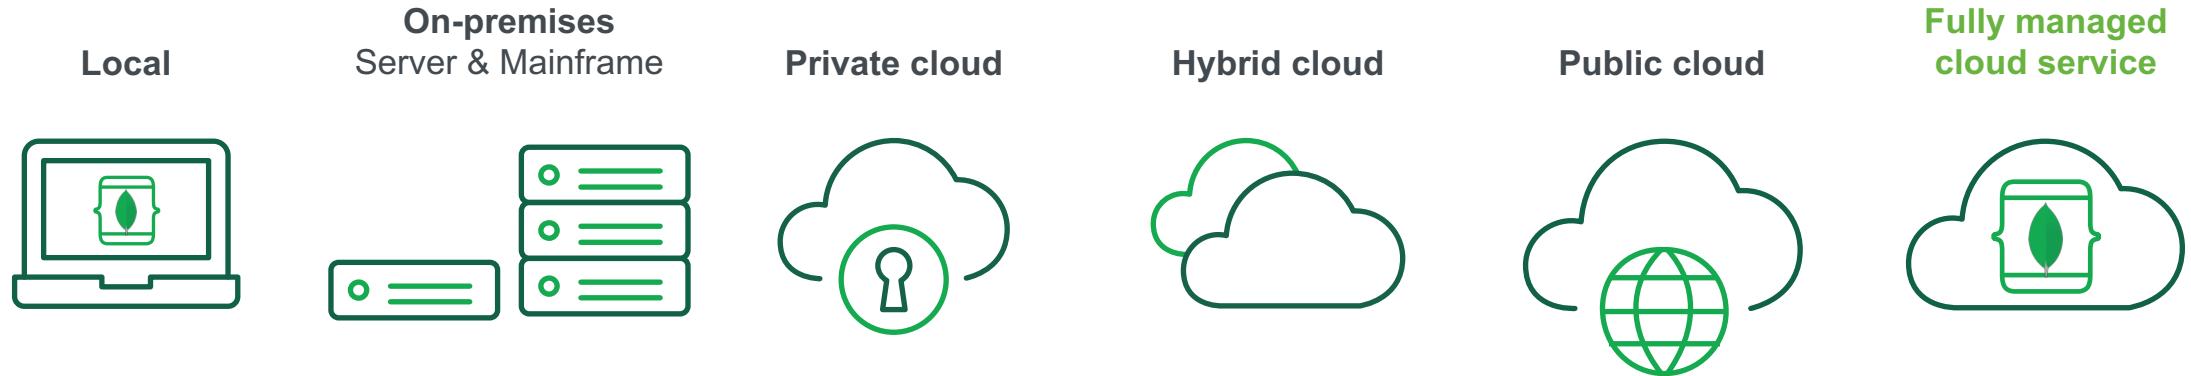


## Sharding

- Automatically scale beyond the constraints of a single node
- Application transparent
- Scale and rebalance incrementally, in real time
- Unlike NoSQL systems that randomly spray data across a cluster, MongoDB exposes multiple data distribution policies to optimize for query patterns and locality

# Run Anywhere

---



- Database that runs the same everywhere
- Leverage the benefits of a multi-cloud strategy
- Global coverage
- Avoid lock-in

**Convenience:** same codebase, same APIs, same tools, wherever you run



# MongoDB Server

---



## mongoDB® Atlas

- The fastest, easiest way to get started with MongoDB.
- <https://www.mongodb.com/cloud/atlas>

Downloadable solutions are available too, but must be self managed.

- Community Edition
- Enterprise Edition



# mongoDB® Atlas : The only *true* multi-cloud database as a service



Google Cloud Platform



## Database that runs the same everywhere

Consistent experience  
across AWS, Azure, and  
GCP

## Coverage in any geography

Deploy in over 50 regions  
worldwide  
  
Create globally distributed  
databases with a few clicks

## Leverage the benefits of a multi-cloud strategy

Exploit the benefits of AWS,  
Azure, or GCP services on  
your data

## Avoid lock-in

Easily migrate data  
between cloud providers

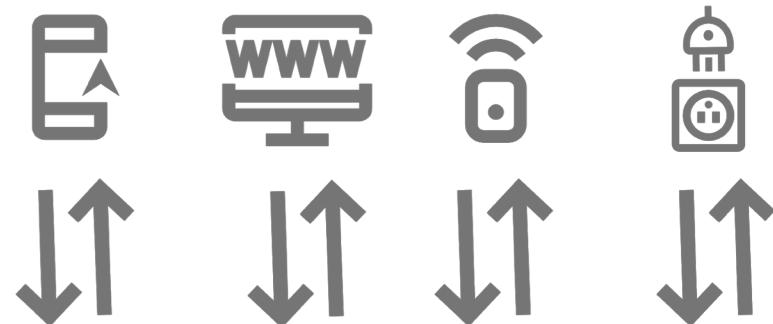
# MongoDB Serverless

---



# mongoDB<sup>®</sup> Stitch

- The serverless platform from MongoDB.
- <https://www.mongodb.com/cloud/stitch>





# MongoDB GUI

---



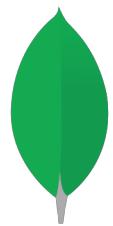
# mongoDB® Compass

- The easiest way to explore and manage your MongoDB Data
- <https://www.mongodb.com/products/compass>



# MongoDB for Analytics

---



## mongoDB® Charts



mongoDB® | Connector for Apache Spark



mongoDB® | Connector for BI



# Setting up MongoDB

9:30 – 10:00

---



# MongoDB 101 Workshop

---

Introduction to MongoDB (30 mins)

## Setting up MongoDB (30 mins)

- Laptop: Standalone Engine & Shell, Compass, MongoDB Driver (Python)
- Data Generator: MGenerateJS
- Cloud: MongoDB Atlas account

## Querying, Indexing and Aggregation (30 mins)

- Using Compass, MongoDB Shell & Python

## Schema Design & Use Cases (30 mins)

- Embedding & Referencing, 1-1, 1-M, M-M, Patterns
- Single View, IoT, Time Series, eCommerce/Content Management

## MongoDB Atlas (60 mins)

- Management, Clusters, Security
- High Availability & Scalability ( Replication & Sharding )
- Managed BI Connector & Charts
- ODBC for BI Connector
- MongoDB Stitch\*



# Downloads & Setup

---

MongoDB Server

<https://www.mongodb.com/download-center/community>

MongoDB Compass - 1.16.4 (stable version)

<https://www.mongodb.com/download-center/compass>

Node.JS and NPM (Node Package Manager)

<https://nodejs.org/en/download/>

<https://www.npmjs.com/get-npm>

MGenerateJS (random data generator)

<https://www.npmjs.com/package/mgeneratejs>

Python / Anaconda

<https://www.python.org/downloads/>

<https://www.anaconda.com/distribution/>

PIP (Package Installer for Python)

<https://pip.pypa.io/en/stable/installing/>

Python MongoDB Driver

<http://api.mongodb.com/python/current/installation.html>

GitHub

<https://desktop.github.com/>

Clone <https://github.com/snarvaez/mdb-101>



# Run MongoDB Server (local)

---

<https://docs.mongodb.com/manual/installation/>

Terminal 1 (Server):

```
./mongod --dbpath ../data
```

Terminal 2 (Shell):

```
./mongo
```

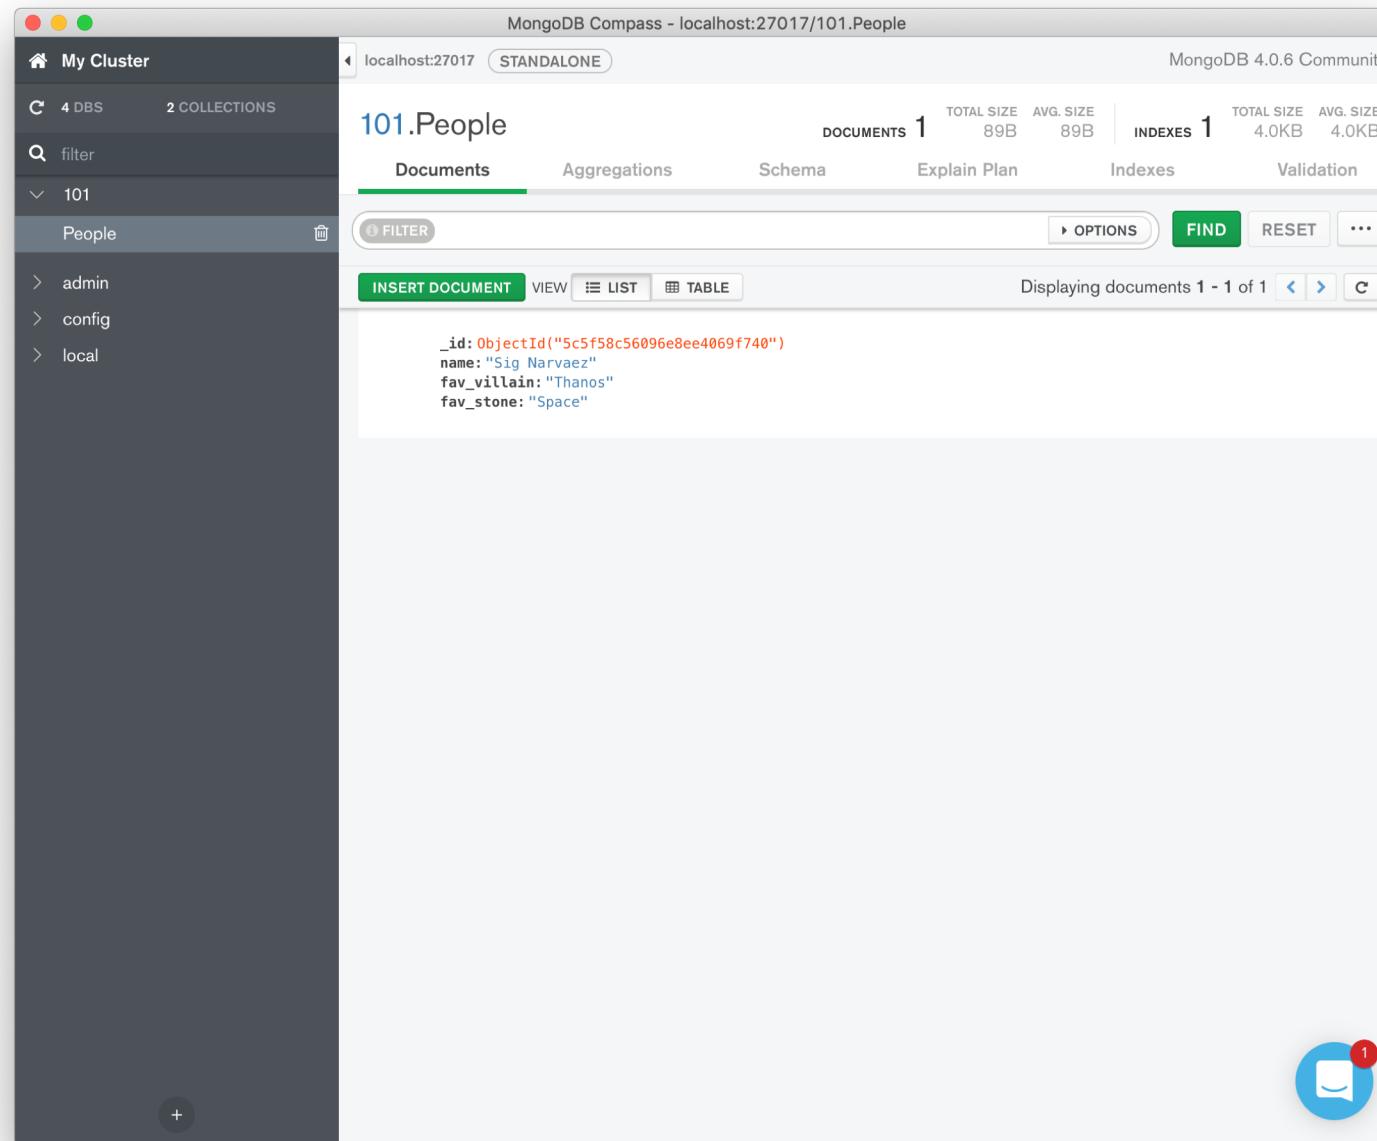
```
> use 101
switched to db 101

> db.People.save({ "name": "Sig Narvaez", "fav_villain": "Thanos" });
WriteResult({ "nInserted" : 1 })

> db.People.count();
1

> db.People.find({ "fav_villain": "Thanos" });
{ "_id" : ObjectId("5c5f53430d71dc70544ef715"), "name" : "Sig Narvaez", "fav_villain" : "Thanos" }
```

# Run MongoDB Compass



The screenshot shows the MongoDB Compass interface running on localhost:27017. The title bar reads "MongoDB Compass - localhost:27017/101.People". The left sidebar shows "My Cluster" with "4 DBS" and "2 COLLECTIONS". Under "101", there are two collections: "People" (selected) and "config". The main panel displays the "101.People" collection with one document listed:

```
_id: ObjectId("5c5f58c56096e8ee4069f740")
name: "Sig Narvaez"
fav_villain: "Thanos"
fav_stone: "Space"
```

Below the document list are tabs for "Documents", "Aggregations", "Schema", "Explain Plan", "Indexes", and "Validation". There are also buttons for "INSERT DOCUMENT", "VIEW", "LIST", and "TABLE". The status bar at the bottom indicates "Displaying documents 1 - 1 of 1".



# Setup MongoDB Driver

---

<https://docs.mongodb.com/ecosystem/drivers/>

<https://docs.mongodb.com/ecosystem/drivers/python/>

<http://api.mongodb.com/python/current/installation.html>

<http://api.mongodb.com/python/current/tutorial.html>

python -m pip install pymongo

git clone <https://github.com/snarvaez/mdb-101>

python mdb-101.py



# Setup MongoDB Data Generator

---

Node Package Manager (NPM) and Node.js

- <https://docs.npmjs.com/downloading-and-installing-node-js-and-npm>

MGenerateJS

- <https://github.com/rueckstiess/mgeneratejs/projects>

```
mgeneratejs -n 1 people_template.json
```

```
mgeneratejs -n 1 people_template.json | ./mongoimport -d 101 -c People
```

```
mgeneratejs -n 10000 people_template.json | ./mongoimport -d 101 -c People
```

Extra time?

- Create variance in your collection(s)
- Change the template (e.g. add/remove fields)



# MongoDB Atlas Account

---

[cloud.mongodb.com](https://cloud.mongodb.com)

New Account

Use Credit Code & Setup Credit Card (optional)

Create M0 Free Cluster

Whitelist current IP

Setup DB User

Insert data with generator

Connect/Browse with UI & Compass

\$75 credit:  
**MDBATRMOUG75**



# Querying, Indexing & Aggregation

10:00 – 10:30

---



# MongoDB 101 Workshop

---

**Introduction to MongoDB** (30 mins)

**Setting up MongoDB** (30 mins)

- Laptop: Standalone Engine & Shell, Compass, MongoDB Driver (Python)
- Data Generator: MGenerateJS
- Cloud: MongoDB Atlas account

**Querying, Indexing and Aggregation** (30 mins)

- Using Compass, MongoDB Shell & Python

**Schema Design & Use Cases** (30 mins)

- Embedding & Referencing, 1-1, 1-M, M-M, Patterns
- Single View, IoT, Time Series, eCommerce/Content Management

**MongoDB Atlas** (60 mins)

- Management, Clusters, Security
- High Availability & Scalability ( Replication & Sharding )
- Managed BI Connector & Charts
- ODBC for BI Connector
- MongoDB Stitch\*

# Use Your Favorite Programming Language

## Drivers

Drivers for most popular programming languages and frameworks



Java



.NET



PHP



Ruby



JavaScript



Perl



Python



```
> db.contact.insert({name: "John Smith", address: "10  
3rd St.", phone: {home: 1234567890, mobile:  
1234568138}})  
>  
> db.contact.findOne()  
{  
    "_id": ObjectId("5106c1c2fc629bfe52792e86"),  
    "name": "John Smith",  
    "address": "10 3rd St.",  
    "phone": {  
        "home": 1234567890,  
        "mobile": 1234568138 }  
}
```

## Shell

Command-line shell for interacting directly with database

mongoDB



# Query Language and Operators

```
// find customers with any claims  
> db.customers.find( {claims: {$exists: true }} )  
  
// find customers matching a regular expression  
> db.customers.find( {last: /^rog*/i } )  
  
// count customers by city  
> db.customers.find( {city: 'Los Angeles'} ).count()
```

## Queries and What They Match

{a: 10}	Docs where <b>a</b> is <b>10</b> , or an array containing the value <b>10</b> .
{a: 10, b: "hello"}	Docs where <b>a</b> is <b>10</b> and <b>b</b> is <b>"hello"</b> .
{a: {\$gt: 10}}	Docs where <b>a</b> is greater than <b>10</b> . Also available: <b>\$lt (&lt;)</b> , <b>\$gte (&gt;=)</b> , <b>\$lte (&lt;=)</b> , and <b>\$ne (!=)</b> .
{a: {\$in: [10, "hello"]}}	Docs where <b>a</b> is either <b>10</b> or <b>"hello"</b> .
{a: {\$all: [10, "hello"]}}	Docs where <b>a</b> is an array containing both <b>10</b> and <b>"hello"</b> .
{"a.b": 10}	Docs where <b>a</b> is an embedded document with <b>b</b> equal to <b>10</b> .
{a: {\$elemMatch: {b: 1, c: 2}}}	Docs where <b>a</b> is an array that contains an element with both <b>b</b> equal to <b>1</b> and <b>c</b> equal to <b>2</b> .
{\$or: [{a: 1}, {b: 2}]}]	Docs where <b>a</b> is <b>1</b> or <b>b</b> is <b>2</b> .
{a: /^m/}	Docs where <b>a</b> begins with the letter <b>m</b> . One can also use the regex operator: <b>{a: {\$regex: "m"}}</b> .
{a: {\$mod: [10, 1]}}	Docs where <b>a mod 10</b> is <b>1</b> .
{a: {\$type: 2}}	Docs where <b>a</b> is a string. See <a href="http://bsonspec.org">bsonspec.org</a> for more.
{ \$text: { \$search: "hello" } }	Docs that contain <b>"hello"</b> on a text search. Requires a text index.

# Secondary Indexes

```
// Create index on policyId (ascending)
> db.policies.createIndex( { policyId : 1 } )
```

```
// Create index on policyId (descending)
> db.policies.createIndex( { policyId : -1 } )
```

```
// Create index on arrays for values on the "types" field – multi key index
> db.policies.createIndex( { types : 1 } ) // { types: ["home", "auto", "life"] }
```

```
// Compound index on the policyId and type fields
> db.policies.createIndex( { policyId : 1, type : 1 } )
```

```
// Compound index supports both the below queries
> db.policies.find( { policyId : 123, type : 'life' } )
> db.policies.find( { policyId : 123 } )
```

```
// Text index on the comments field
> db.policies.createIndex( {comments: "text"} )
```

```
// Search comments and also score text search results
> db.policies.find ({$text: { $search: "payment" }}, { score: { $meta: "textScore" }})
```



# Java - Basic Query API Example

Objective	Code	CLI
<b>Find all contacts with at least one mobile phone</b>	Map <b>expr</b> = new HashMap(); expr.put("phones.type", "mobile");	db.contact.find({"phones.type": "mobile"});
<b>Find contacts with NO phones</b>	Map <b>expr</b> = new HashMap(); Map q1 = new HashMap(); q1.put("\$exists", false); expr.put("phones", q1);	db.contact.find({"phones": {"\$exists": false}});

```
List fetchGeneral(Map expr)
{
    List l = new ArrayList();
    DBObject dbo = new BasicDBObject(expr);
    Cursor c = collection.find(dbo);
    while (c.hasNext()) {
        Map m = c.next();
        showMap(m);
        l.add(m);
    }
    return l;
}

// collection.insert(Map m)
// collection.update(Map filter, Map repl)
// collection.remove(Map filter)
```

```
{
    customer_id : 1,
    first_name : "Mark",
    last_name : "Smith",
    city : "San Francisco",
    phones: [ {
        number : "1-212-777-1212", dnc : true,
        type : "home"
    }, {
        number : "1-212-777-1213",
        type : "cell"
    } ] }
```



# C# Examples

```
// To directly connect to a single MongoDB server
// (this will not auto-discover the primary even if it's a member of a replica set)
var client = new MongoClient();

// or use a connection string
var client = new MongoClient("mongodb://localhost:27017");

// or, to connect to a replica set, with auto-discovery of the primary, supply a seed list of
var client = new MongoClient("mongodb://localhost:27017,localhost:27018,localhost:27019");

var database = client.GetDatabase("foo");

var collection = database.GetCollection<BsonDocument>("bar");

var count = await collection.CountAsync(new BsonDocument());

Console.WriteLine(count);

var document = await collection.Find(new BsonDocument()).FirstOrDefaultAsync();
Console.WriteLine(document.ToString());
```

# Python Example

```
import pymongo

myData = {
    "name": "jane",
    "id": "K2",
    # no title? No problem
    "hireDate": datetime.date(2011, 11, 1),
    "phones": [
        { "type": "work",
          "number": "1-800-555-1212"
        },
        { "type": "home",
          "number": "1-866-444-3131"
        }
    ]
}
coll.insert(myData)
print coll.find_one({"bizKey": "K2" })

expr = { "$or": [ {"phones": { "$exists": false }}, {"name": "jane"} ] }
for c in coll.find(expr):
    print [ k.upper() for k in sorted(c.keys()) ]
```

- **Easy and fast to create scripts due to "fidelity-parity" of MongoDB map data and python (and perl, ruby, and Javascript) structures**
- Data types and structure in scripts are exactly the same as that read and written in Java and C++



# Rich query functionality

---

Expressive Queries	<ul style="list-style-type: none"><li>Find anyone with phone # “1-212...”</li><li>Check if the person with number “555...” is on the “do not call” list</li></ul>
Geospatial	<ul style="list-style-type: none"><li>Find the best offer for the customer at geo coordinates of 42nd St. and 6th Ave</li></ul>
Text Search	<ul style="list-style-type: none"><li>Find all tweets that mention the firm within the last 2 days</li></ul>
Aggregation	<ul style="list-style-type: none"><li>Count and sort number of customers by city, compute min, max, and average spend</li></ul>
Native Binary JSON Support	<ul style="list-style-type: none"><li>Add an additional phone number to Mark Smith’s record without rewriting the document</li><li>Update just 2 phone numbers out of 10</li><li>Sort on the modified date</li></ul>
JOIN (\$lookup)	<ul style="list-style-type: none"><li>Query for all San Francisco residences, lookup their transactions, and sum the amount by person</li></ul>
Graph Queries (\$graphLookup)	<ul style="list-style-type: none"><li>Query for all people within 3 degrees of separation from Mark</li></ul>

## MongoDB

```
{   customer_id : 1,  
    first_name : "Mark",  
    last_name : "Smith",  
    city : "San Francisco",  
    phones: [      {  
        number : "1-212-777-1212",  
        type : "work"  
    },  
    {  
        number : "1-212-777-1213",  
        type : "cell"  
    }]  
.... . . .
```



# Indexes in MongoDB

---

*Indexes support the efficient execution of queries in MongoDB.*

# Can still query w/o index → collscan

blog.users

Documents

FILTER password

COLLSCAN

Query Performance Summary

- Documents Returned: 1
- Index Keys Examined: 0
- Documents Examined: 100000
- Actual Query Execution Time (ms): 38
- Sorted in Memory: no
- No index available for this query.

nReturned: 1

Documents Examined: 100000

DETAILS

The screenshot shows the MongoDB Compass interface. At the top, there's a decorative header with colored bars. Below it, the database 'blog' and collection 'users' are selected. The top bar displays statistics: DOCUMENTS 100.0k, TOTAL SIZE 9.0MB, AVG. SIZE 94B, INDEXES 1, TOTAL SIZE 1.0MB, and AVG. SIZE 1.0MB. A large red box highlights the 'COLLSCAN' status in the top bar and the 'No index available for this query.' message in the 'Query Performance Summary' section. The 'Query Performance Summary' section also lists other metrics: Documents Returned: 1, Index Keys Examined: 0, and Documents Examined: 100000. The bottom part of the interface shows a table with 'nReturned: 1' and a 'DETAILS' button, which is also highlighted by a red box.

# Index Creation in Compass

---

Create Index

Choose an index name

Configure the index definition

Select a field name  Select a type

karma  
lang  
**password**  
registered

Create unique index  
 Create TTL  
 seconds

Partial Filter Expression

# Search with Index

blog.users

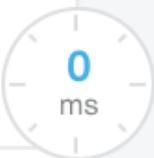
Documents Aggregations Schema Explain

FILTER { password: "lXzCyCotsf"}

## FETCH

nReturned: 1

Execution Time:



DETAILS

## Query Performance Summary

i Documents Returned: 1

i Index Keys Examined: 1

i Documents Examined: 1

i Actual Query Execution Time (ms): 0

i Sorted in Memory: no

i Query used the following index:

password

## IXSCAN

nReturned: 1

Execution Time: 0 ms

Index Name: password

Multi Key Index: no

DETAILS

Multi Key Index: no

DETAILS



# Index Types

---

Single Field

Compound Field

Multikey

Geospatial

Text

Hashed

# Index Properties

Unique

Partial

Sparse

TTL

<https://docs.mongodb.com/manual/indexes>

# Aggregation Pipeline

---

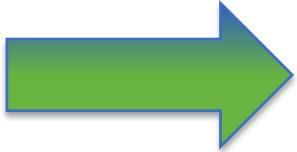


# PIPELINE

---

\*nix command line pipe

```
ps ax | grep mongod| head 1
```



# PIPELINE

---

MongoDB document pipeline

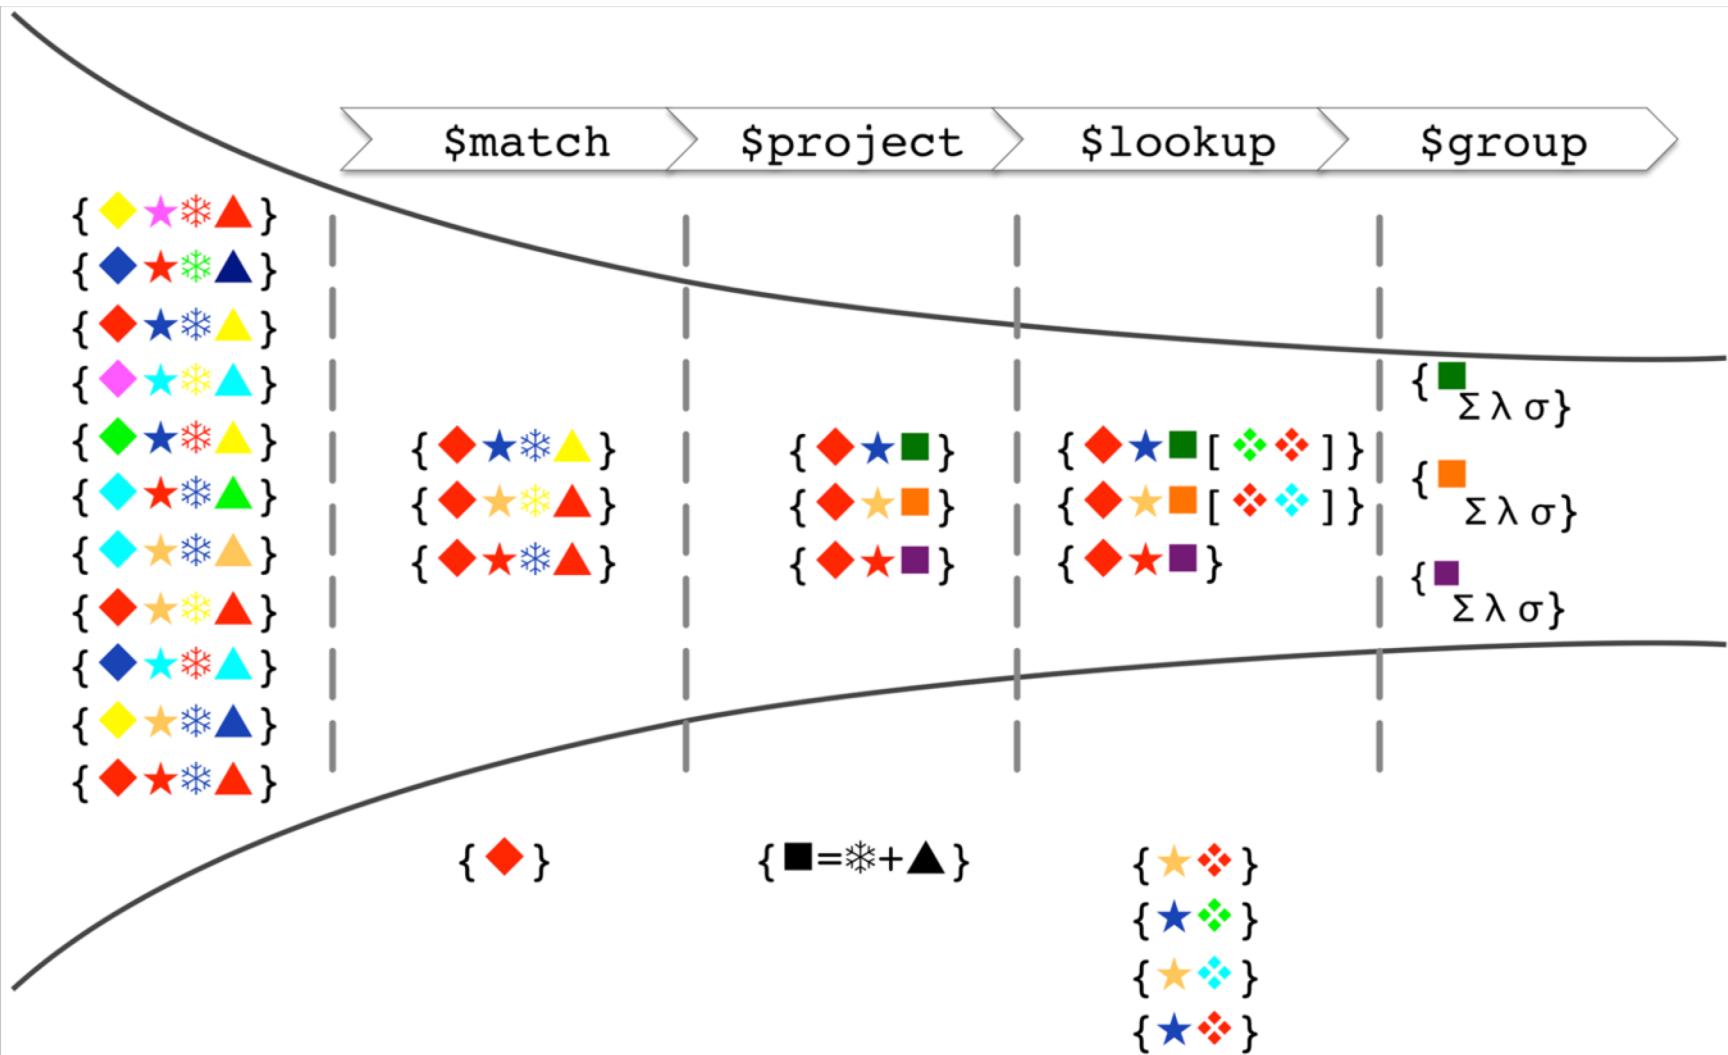
\$match | \$group | \$sort



Input stream {}  
{} {} {}

→ Result {} {}  
...

# Aggregation



```
{ title: "The Great Gatsby",  
  language: "English",  
  subjects: [  
    "Long Island",  
    "New York",  
    "1920s" ],  
  { title: "War and Peace",  
    language: "Russian",  
    subjects: [  
      "Russia",  
      "War of 1812",  
      "Napoleon" ],  
  { title: "Open City",  
    language: "English",  
    subjects: [  
      "New York",  
      "Harlem" ] }  
  {"$match": {"language": "English"}},  
  {"$unwind": "$subjects"}  
}  
  
{title: "The Great Gatsby",  
  language: "English",  
  subjects: [  
    "Long Island",  
    "New York",  
    "1920s" ],  
  {title: "Open City",  
    language: "English",  
    subjects: [  
      "New York",  
      "Harlem" ] }  
}  
  
{title: "The Great Gatsby",  
  language: "English",  
  subjects: [  
    "Long Island",  
    "New York",  
    "1920s" ],  
  {title: "Open City",  
    language: "English",  
    subjects: "Harlem"}  
}  
  
{ _id: "Long Island",  
  { _id: "New York",  
    { _id: "1920s",  
      count: 1 },  
    { _id:  
      "Harlem",  
      count: 1 },  
  }  
}  
  
{ _id: "New York",  
  { count: 1 },  
  { _id: "1920s",  
    count: 1 },  
  { _id: "Harlem",  
    count: 1 },  
}  
  
{"$sort": {"count": -1}}  
{"$limit": 3}  
{"$project": {...}}
```

\$match

\$unwind

\$group

\$sort

\$limit

\$project

\$skip

# Aggregation Pipeline Builder in Compass

---

# Pipeline Builder

The screenshot shows the MongoDB Pipeline Builder interface for the `blog.users` collection. The top navigation bar includes tabs for `Documents`, `Aggregations` (which is highlighted with a red circle), `Schema`, `Explain Plan`, `Indexes`, and `Validation`. Below the tabs are buttons for `COMMENT MODE`, `SAMPLE MODE`, and `AUTO PREVIEW`. A message indicates there are 100000 documents in the collection.

In the main area, a message says "Select an operator to construct expressions used in the aggregation pipeline stages." Below this, a `$match` stage is defined with the query `{karma: {$gte: 250}}`. The output section shows a sample of 20 documents from the collection.

`_id: "USER_10"`  
`password: "ijwgrASaAl"`  
`karma: 327`  
`registered: 2019-01-04 12:58:59.315`  
`lang: "EN"`

`_id: "USER_10"`  
`password: "ijwgrASaAl"`  
`karma: 16`  
`registered: 2019-01-04 12:58:59.315`  
`lang: "EN"`

`_id: "USER_10"`  
`password: "ijwgrASaAl"`  
`karma: 49`  
`registered: 2019-01-04 04:58:59.315`  
`lang: "EN"`

**Output after \$match stage (Sample of 20 documents)**

`1 *  
2 * query - The query in MQL.  
3 */  
4 {  
5 karma: {$gte: 250}  
6 }`

**ADD STAGE**

# Pipeline Builder

---

The screenshot shows the MongoDB Aggregations interface for the `blog.users` collection. The interface has three tabs: `Documents`, `Aggregations` (which is selected), and `Schema`. Below the tabs is a toolbar with buttons for navigating between documents, saving the pipeline, and more. A dropdown menu is open over the `SAVE PIPELINE` button, showing options: `Export To Language` (highlighted in blue), `Clone Pipeline`, `New Pipeline`, and `New Pipeline From Text`.

# Pipeline Builder Language Exporter

## Export Pipeline To Language

My Pipeline:

```
1 [{  
2   $match: {  
3     karma: {  
4       $gte: 250  
5     }  
6   }, {  
7     $sort: {  
8       registered: 1,  
9       karma: -1  
10    }  
11  }]  
12 }]
```



Export Pipeline To:

PYTHON 3

JAVA  
NODE  
C#  
**PYTHON 3**

```
1 [{  
2   '$mat  
3   }  
4   }, {  
5     '$sort': {  
6       'registered': 1,  
7       'karma': -1  
8     }  
9   }]  
10 }  
11 ]  
12 }  
13 }  
14 }
```



Include Import Statements

CLOSE

mongoDB®



# Hands on!

---

## **Customer Single View Template**

```
mgeneratejs -n 100000 CustomerSingleView.json | ./mongoimport -d 101 -c Customers
```

# Complex Query

Show all insurance customers who are Female, born in 1990, live in Utah and have at least one Life insurance policy for which they are registered as a Smoker

FILTER: {'address.state': 'UT', policies: {\$elemMatch: {'policyType': 'life', 'insured\_person.smoking': true}}, yob: {\$gte: ISODate('1990-01-01'), \$lte: ISODate('1990-12-12')}, 'gender': 'Female'}

PROJECT: {\_id:0, firstname:1, lastname:1, yob: 1}

**test.customers**

DOCUMENTS	311.0k	TOTAL SIZE	AVG. SIZE	INDEXES	1	TOTAL SIZE	AVG. SIZE
380.3MB	1.3KB	2.8MB	2.8MB				

**Documents**    **Aggregations**    **Schema**    **Explain Plan**    **Indexes**    **Validation**

**FILTER** : {\$elemMatch: {'policyType': 'life', 'insured\_person.smoking': true}}, yob: {\$gte: ISODate('1990-01-01'), \$lte: ISODate('1990-12-12')}, gender: 'Female'    **OPTIONS**    **FIND**    **RESET**    **...**

**PROJECT** : {\_id:0, firstname:1, lastname:1, yob: 1}

**SORT**

**COLLATION**

**Skip** 0    **Limit** 0

**VIEW**    **LIST**    **TABLE**

Displaying documents 1 - 20 of 55

mongoDB®



# Explain Plan

Documents Aggregations Schema Explain Plan Indexes Validation

**FILTER** \$elemMatch: {'policyType': 'life', 'insured\_person.smoking': true}, yob: {\$gte: ISODate('1990-01-01'), \$l OPTIONS

**PROJECT** {\_id:0, firstname:1, lastname:1, yob: 1}

**SORT**

**COLLATION**

**SKIP** 0 **LIMIT** 0

**EXPLAIN** **RESET** ...

VIEW DETAILS AS **VISUAL TREE** RAW JSON

Documents Returned: 55

Index Keys Examined: 0

Documents Examined: 1000000

Actual Query Execution Time (ms): 6925

Sorted in Memory: no

No index available for this query.

**PROJECTION**  
nReturned: 55 Execution Time: 0 ms  
Transform by:  
{ "\_id":0,"firstname":1,"lastname":1,"yob":1}  
**DETAILS**

**COLLSCAN**  
nReturned: 55 Execution Time: 6895 ms  
Documents Examined: 1000000

# Indexes

Create Index

Choose an index name

GenderLifeInsuranceSmokersDOBByStateIndex

Configure the index definition

address.state	1 (asc)	-
policies.policyType	1 (asc)	-
policies.insured_person.smoki...	1 (asc)	-
yob	1 (asc)	-
gender	1 (asc)	-

**ADD ANOTHER FIELD**

▼ Options

Build index in the background

Create unique index

Create TTL

Partial Filter Expression

{ }

CANCEL CREATE

# Efficient Query

Test	# Docs Examined	# Docs Returned	Exec Time (ms)
Rich query without index	1000000	~50	6000
Rich query with index	~50	~50	1

DOCUMENTS 311.0k
TOTAL SIZE 380.3MB
AVG. SIZE 1.3KB
INDEXES 1
TOTAL SIZE 2.8MB
AVG. SIZE 2.8MB

**test.customers**

- Documents
- Aggregations
- Schema
- Explain Plan**
- Indexes
- Validation

**FILTER** `!elemMatch: { 'policyType': 'life', 'insured_person.smoking': true}}, yob: {$gte: ISODate('1990-01-01'), $l`
▼ OPTIONS

**PROJECT** `{ '_id': 0, 'firstname': 1, 'lastname': 1, 'yob': 1 }`

**SORT**

**COLLATION**

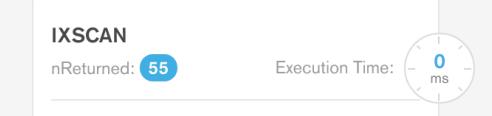
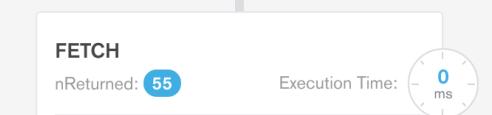
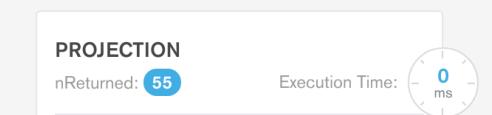
SKIP 0
LIMIT 0

VIEW DETAILS AS
VISUAL TREE
RAW JSON

**Query Performance Summary**

<b>Documents Returned:</b> 55	<b>Actual Query Execution Time (ms):</b> 0
<b>Index Keys Examined:</b> 129	<b>Sorted in Memory:</b> no
<b>Documents Examined:</b> 55	<b>Query used the following index:</b>

address.state ↑ policies.policyType ↑ policies.i...



Multi Key Index: yes



# Aggregation

---

*Show all Sum of premiums for policies due in May 2018 by Policy Type*

```
1 // Sum of premiums for policies due in May 2018 by Policy Type
2 db.Customers.aggregate([
3   { "$match" : { 'policies.nextRenewalDt': {
4     "$gte": new Date('2018-05-01'),
5     "$lte": new Date('2018-05-30') }}},
6
7   { "$unwind" : "$policies"},  

8
9   { "$group" : { "_id" : "$policies.policyType",
10     "renewalTotal" : { "$sum" : "$policies.premiumYear" }}}
11 ]);
```

# Aggregation

Show all Sum of premiums for policies due in May 2018 by Policy Type

MongoDB Compass - localhost:27017/101.Customers      MongoDB 4.0.6 Community

101.Customers      DOCUMENTS 100.0k      TOTAL SIZE 129.5MB      AVG. SIZE 1.3KB      INDEXES 2      TOTAL SIZE 6.3MB      AVG. SIZE 3.1MB

Documents    Aggregations    Schema    Explain Plan    Indexes    Validation

Premiums Due May 2018 | SAVE PIPELINE | ...

Preview of Documents in the Collection

\$match    Output after \$match stage (Sample of 20 documents)

```
1- {
2-   "policies.nextRenewalDt": {
3-     "$gte": ISODate("2018-05-01"),
4-     "$lte": ISODate("2018-05-30")
5-   }
6 }
```

\$unwind    Output after \$unwind stage (Sample of 20 documents)

```
1- {
2-   path: "$policies"
3 }
```

\$group    Output after \$group stage (Sample of 3 documents)

```
1- {
2-   "_id" : "Spolicies.policyType",
3-   "renewalTotal" : { "$sum" : "$policies.premium"
4 }
```

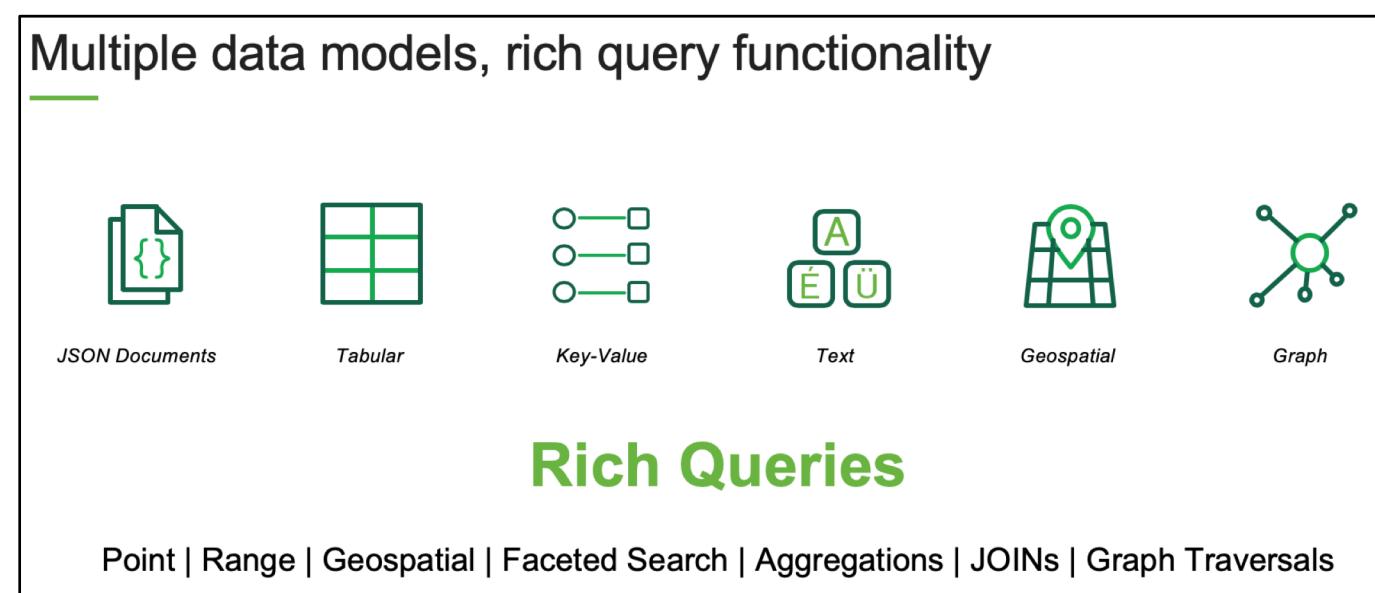
# Repeat via MongoDB shell & Python

Time Check ??

QueryIndexAgg.js

QueryIndexAgg.py

Multi-Model.js





# Schema Design & Use Cases

10:30 – 11:00

---



# MongoDB 101 Workshop

---

**Introduction to MongoDB** (30 mins)

**Setting up MongoDB** (30 mins)

- Laptop: Standalone Engine & Shell, Compass, MongoDB Driver (Python)
- Data Generator: MGenerateJS
- Cloud: MongoDB Atlas account

**Querying, Indexing and Aggregation** (30 mins)

- Using Compass, MongoDB Shell & Python

**Schema Design & Use Cases** (30 mins)

- Embedding & Referencing, 1-1, 1-M, M-M, Patterns
- Single View, IoT, Time Series, eCommerce/Content Management

**MongoDB Atlas** (60 mins)

- Management, Clusters, Security
- High Availability & Scalability ( Replication & Sharding )
- Managed BI Connector & Charts
- ODBC for BI Connector
- MongoDB Stitch\*

# Tabular/Relational

# vs. Object Oriented





# Documents are FLEXIBLE

---

Different Documents in the same ProductsCatalog collection in MongoDB

Polymorphic Schema – Aligns with OOP principles

Car Paint products

```
{  
  sku: 'PAINTZXC123',  
  product_name: 'Metallic Paint',  
  colors: ['Red', 'Green'],  
  size_gallons: [5, 10]  
}
```

Car T-Shirt products

```
{  
  sku: 'TSHRTASD43546',  
  product_name: 'T-shirt',  
  size: ['S', 'M', 'L', 'XL'],  
  colors: ['Heather Gray' ... ],  
  material: '100% cotton',  
  wash: 'cold',  
  dry: 'tumble dry low'  
}
```

Car Wheels products

```
{  
  sku: 'WHEELBVCX6543',  
  product_name: '19" 5-spoke',  
  material: 'aluminum alloy',  
  color: 'silver',  
  frame_material: 'aluminum',  
  package_height: '20.5x32.9x55',  
  weight_lbs: 5.15,  
  wheel_size_in: 19  
}
```

# Flexible Data Governance

```
db.runCommand({collMod: "ProductCatalog",
  validator: {
    "$and": [
      {"sku": {"$type": "string"}},
      {"product_name": {"$type": "string"}}
    ]
  },
  "validationLevel": "off | moderate | strict",
  "validationAction": "error | warn"
});
```

*All products must have  
an SKU and Name*

# Relational – schema validation

```
db.runCommand({collMod: "orders",
  validator : {
    "$and" : [
      {"orderID" : {"$type" : "string"}},
      {"customer.customerID" : {"$type": "string"}},
      {"items.0" : {"$exists" : true}}
    ],
    "validationLevel": "strict",
    "validationAction": "error"
  }
});
```

Rule

*"An order must have at least one item"*

Not possible with SQL DDL

**Level:** off | moderate | strict  
**Action:** error | warn

```
db.orders.insert({
  "orderID": "order1",
  "customer": {"customerID": "customer1"},
  "items": []
});
WriteResult({
  "nInserted" : 0,
  "writeError" : {
    "code" : 121,
    "errmsg" : "Document failed validation"
  }
})
```

Schema on-write



# Schema Design

---



# It's all about the data

---

## Questions to keep in mind

Data access patterns?      Number of Reads vs. Updates?

Expected size of a Document?

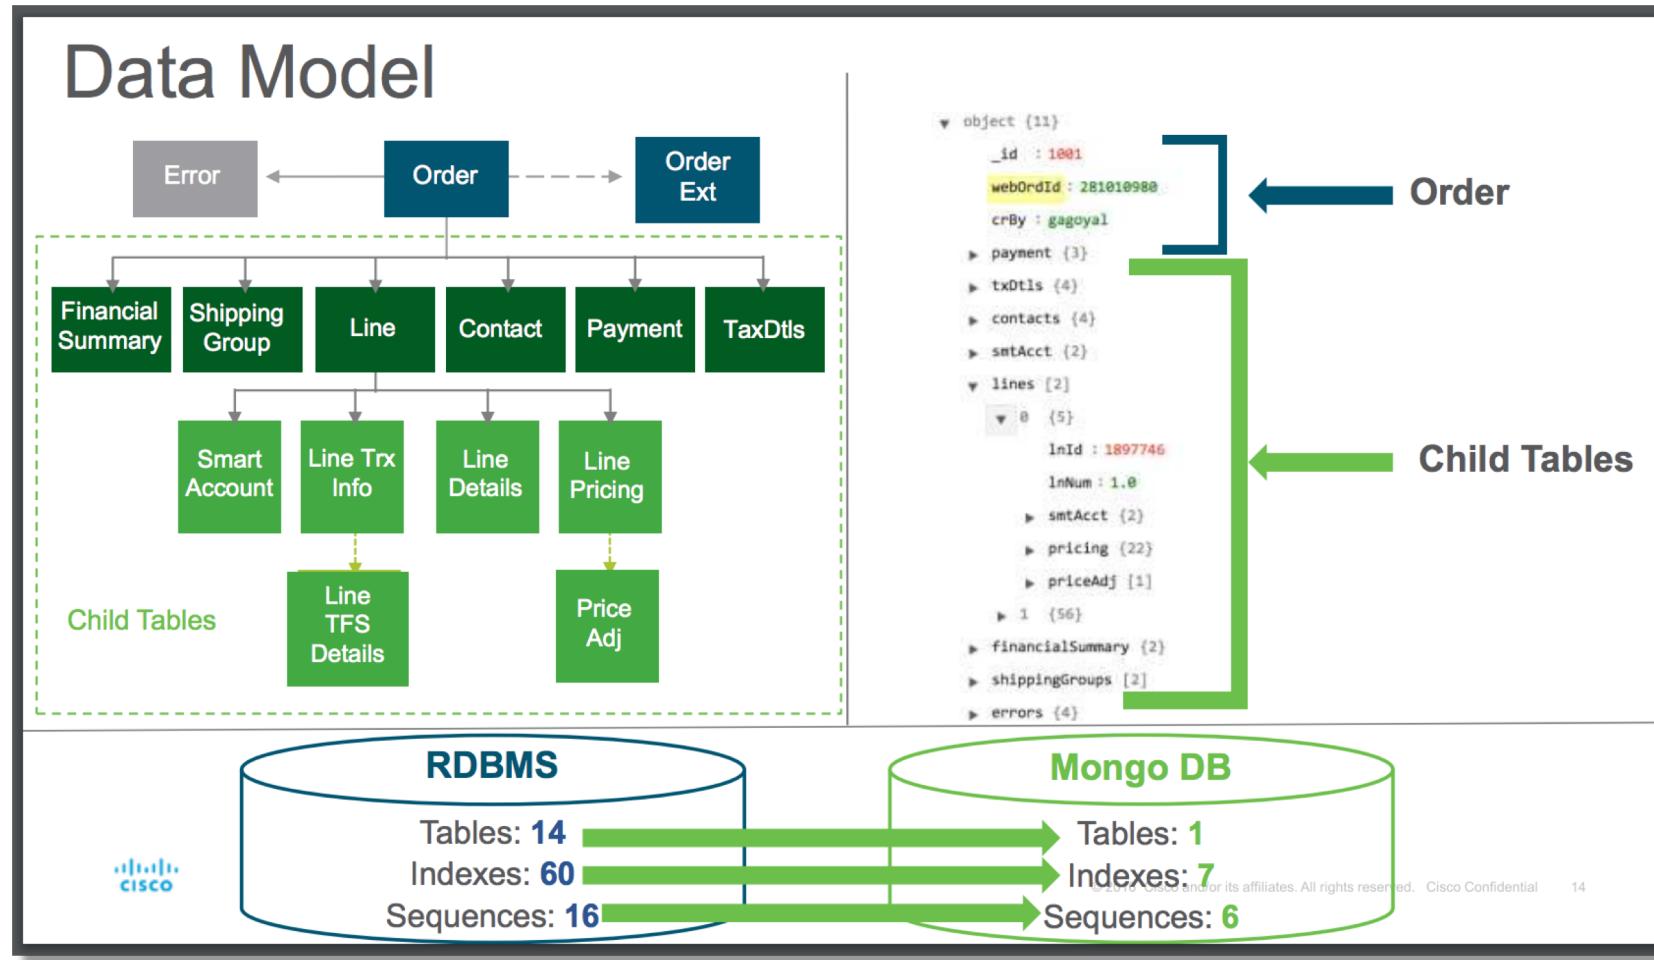
*Data that works together, lives together*

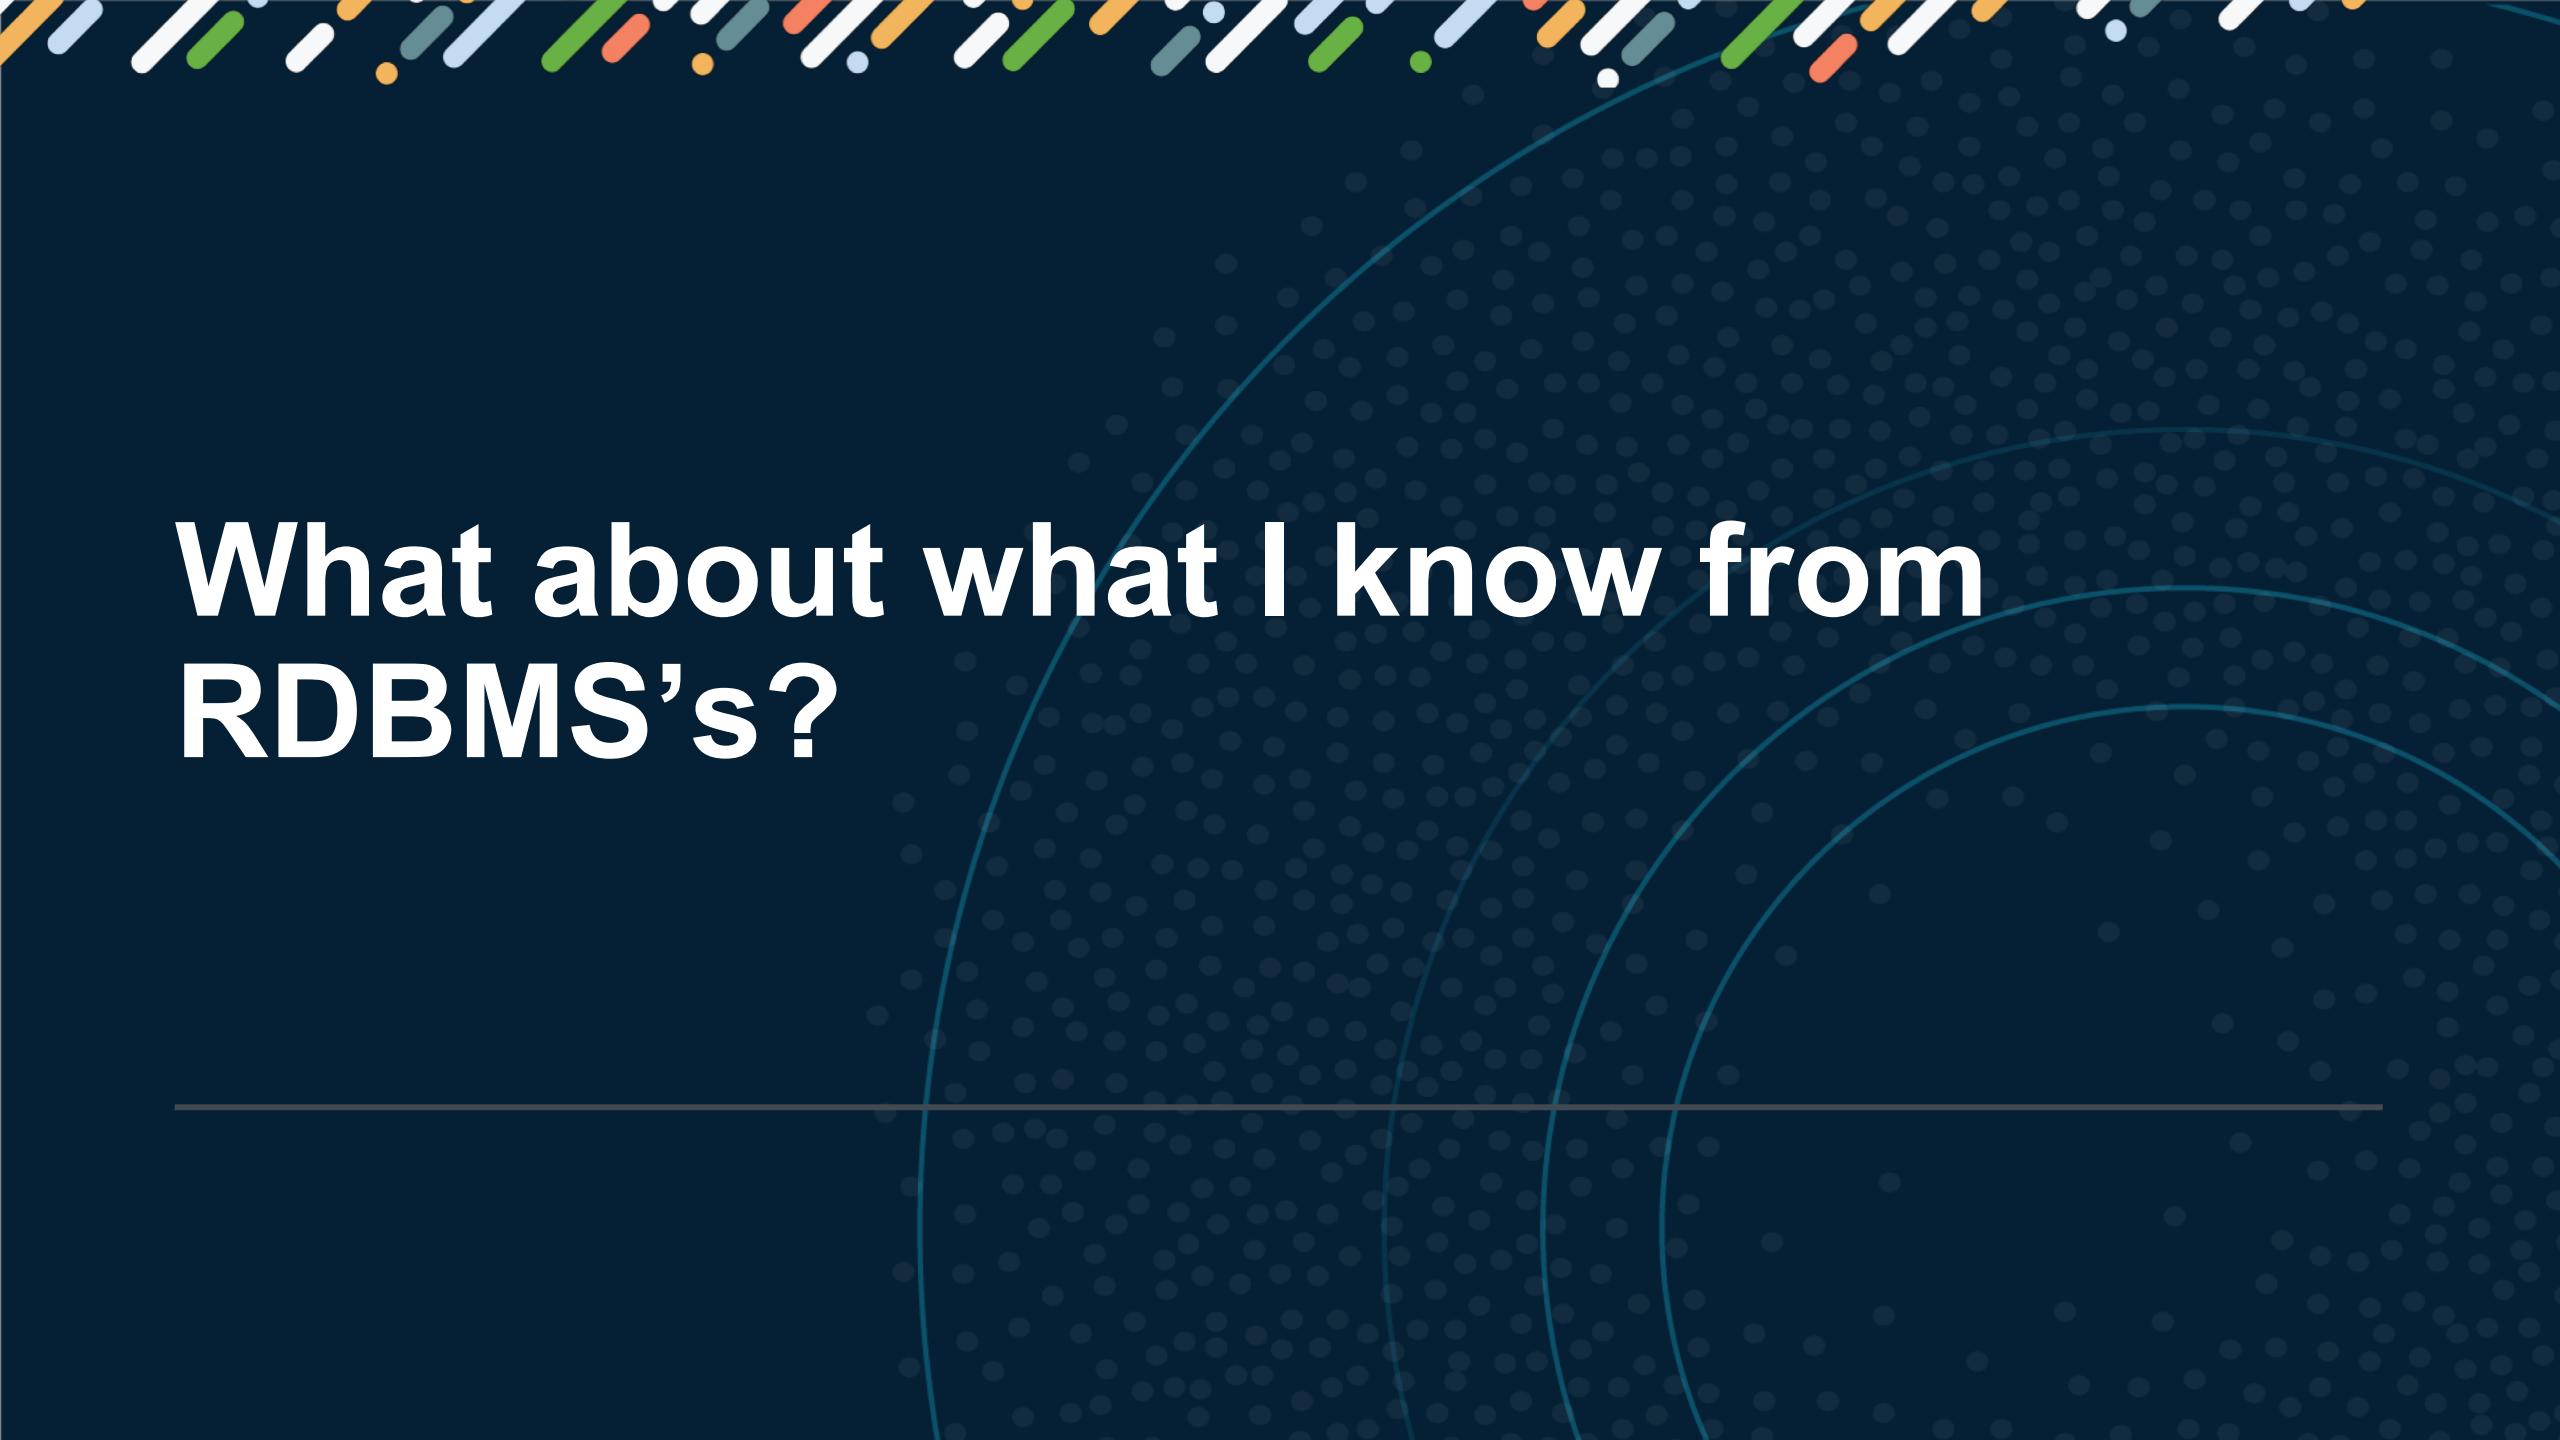
## Why Schema Design is important

Almost all performance issues are related to Schema Design

# CISCO's \$40B eCommerce Platform

14 SQL tables → 1 MongoDB Collection  
60 Indexes → 7 Indexes

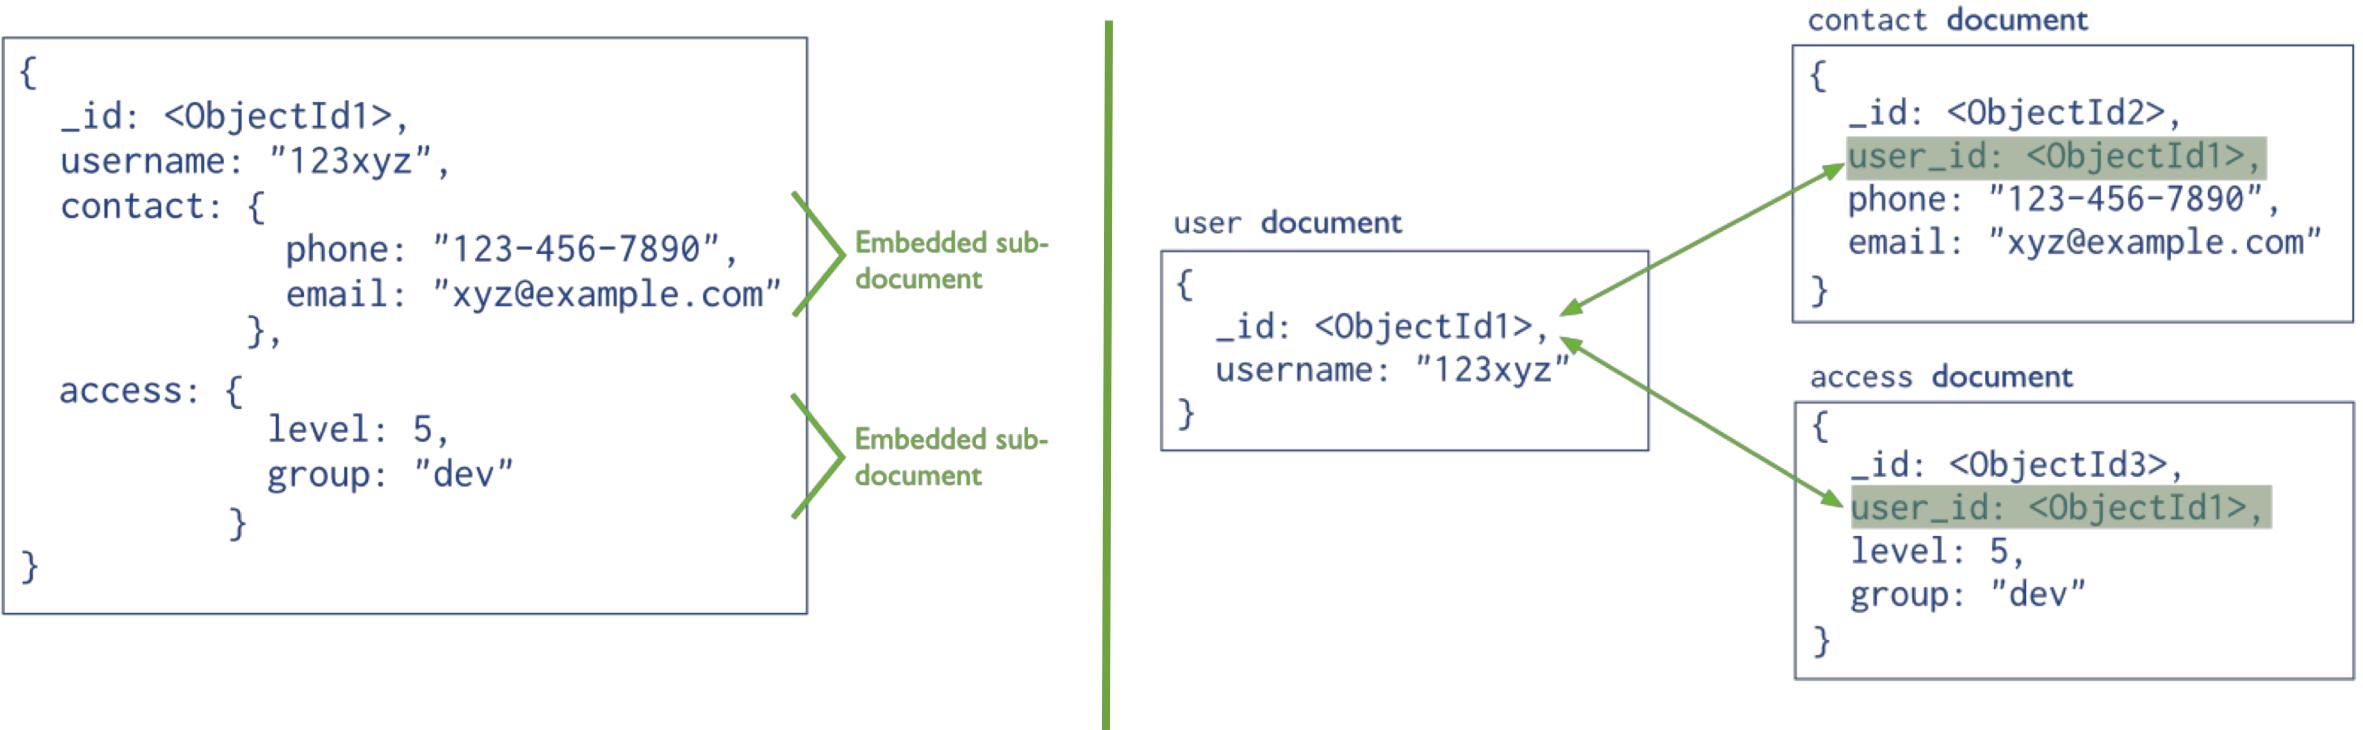




What about what I know from  
RDBMS's?

---

# EMBEDDING & REFERENCING



<https://docs.mongodb.com/manual/core/data-modeling-introduction/>



# Healthcare use case

---





## Medical Procedure collection

1-1

```
{  
    "_id": 333,  
    "date": "2003-02-09T05:00:00",  
    "hospital": "County Hills",  
    "patient": "John Doe",  
    "physician": "Stephen Smith",  
    "procedure": "Glucose",  
    "result": {  
        "value": 97,  
        "measurement": "mg/dl"  
    }  
}
```

Embed – weak entity

# Modeled in 2 possible ways

**1 - M**

## Embed

**Patients**

```
{  
  _id: 2,  
  first: "Joe",  
  last: "Patient",  
  addr: { ...},  
  procedures: [  
    {  
      id: 12345,  
      date: 2015-02-15,  
      type: "Cat scan",  
      ...},  
    {  
      id: 12346,  
      date: 2015-02-15,  
      type: "blood test",  
      ...}  
  ]  
}
```

**Patients**

```
{  
  _id: 2,  
  first: "Joe",  
  last: "Patient",  
  addr: { ...},  
  procedures: [12345, 12346]  
}
```

## Reference

**Procedures**

```
{  
  _id: 12345,  
  date: 2015-02-15,  
  type: "Cat scan",  
  ...}  
{  
  _id: 12346,  
  date: 2015-02-15,  
  type: "blood test",  
  ...}  
}
```

M-M



No Join Tables  
Use **arrays** instead

# Embedding Physicians in Hospitals collection

M-M

```
{  
  _id: 1,  
  name: "Oak Valley Hospital",  
  city: "New York",  
  beds: 131,  
  physicians: [  
    {  
      id: 12345,  
      name: "Joe Doctor",  
      address: {...},  
      ...},  
    {  
      id: 12346,  
      name: "Mary Well",  
      address: {...},  
      ...}  
  ]  
}
```

Data Duplication

...  
is ok!

```
{  
  _id: 2,  
  name: "Plainmont Hospital",  
  city: "Omaha",  
  beds: 85,  
  physicians: [  
    {  
      id: 63633,  
      name: "Harold Green",  
      address: {...},  
      ...},  
    {  
      id: 12345,  
      name: "Joe Doctor",  
      address: {...},  
      ...}  
  ]  
}
```

# Referencing

**M-M**

## Hospitals

```
{  
  _id: 1,  
  name: "Oak Valley Hospital",  
  city: "New York",  
  beds: 131,  
  physicians: [12345, 12346]  
}
```

```
{  
  _id: 2,  
  name: "Plainmont Hospital",  
  city: "Omaha",  
  beds: 85,  
  physicians: [63633, 12345]  
}
```

## Physicians

```
{  
  id: 63633,  
  name: "Harold Green",  
  hospitals: [1,2],  
  ...}  
}
```

```
{  
  id: 12345,  
  name: "Joe Doctor",  
  hospitals: [1],  
  ...}  
}
```

```
{  
  id: 12346,  
  name: "Mary Well",  
  hospitals: [1,2],  
  ...}  
}
```



# MongoDB

## Schema Design Patterns

---

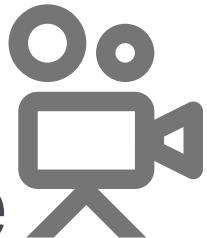
# Full talk from MDBW17

Home » MongoDB World 2017: Session Recordings » Advanced Schema Design Patterns - Daniel Coupal



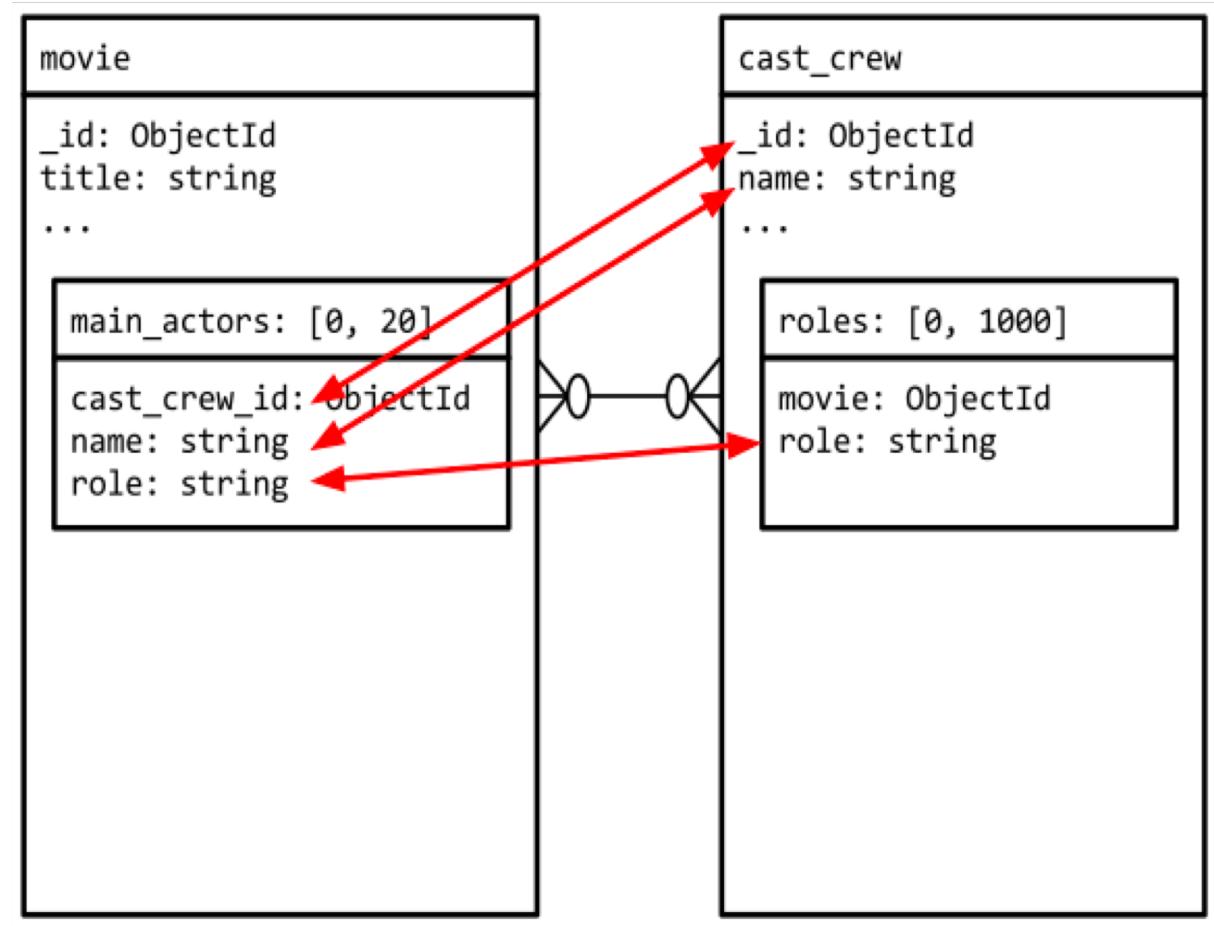
**Advanced Schema Design Patterns -  
Daniel Coupal**

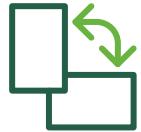
<https://explore.mongodb.com>



# SUBSET Pattern – Entertainment Use Case

- You want to display dependent information, *however only part of it*
- The rest of the data is fetched only if needed
- Examples:
  - The Cast of a Movie
  - Last 10 Movies an Actor has starred in





# Easy: MongoDB Multi-Document ACID Transactions



## Just like relational transactions

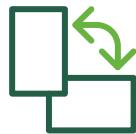
- Multi-statement, familiar relational syntax
- Easy to add to any application
- Multiple documents in 1 or many collections and databases

## ACID guarantees

- Snapshot isolation, all or nothing execution
- No performance impact for non-transactional operations

## Schedule

- MongoDB 4.0: replica set
- MongoDB 4.2: extended to sharded clusters



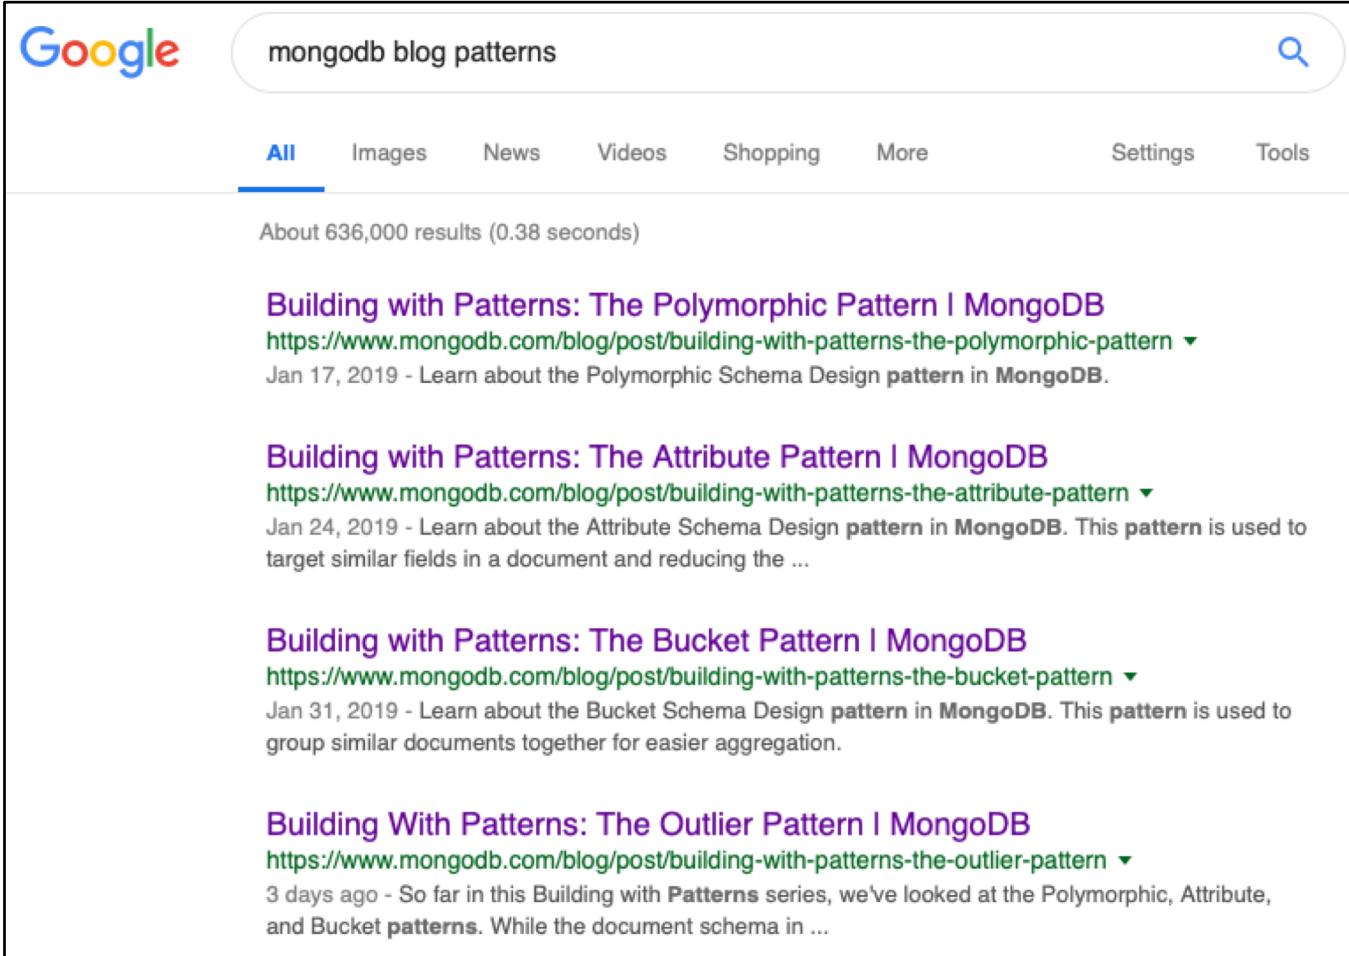
## Syntax

```
with client.start_session() as s:  
    s.start_transaction()  
    collection_one.insert_one(doc_one, session=s)  
    collection_two.insert_one(doc_two, session=s)  
    s.commit_transaction()
```

### Natural for developers

- Idiomatic to the programming language
- Familiar to relational developers
- Simple

# Patterns Blog Posts



A screenshot of a Google search results page. The search query "mongodb blog patterns" is entered in the search bar. The "All" tab is selected, showing approximately 636,000 results found in 0.38 seconds. The results list four blog posts from MongoDB's official blog:

- Building with Patterns: The Polymorphic Pattern | MongoDB**  
<https://www.mongodb.com/blog/post/building-with-patterns-the-polymorphic-pattern> ▾  
Jan 17, 2019 - Learn about the Polymorphic Schema Design pattern in MongoDB.
- Building with Patterns: The Attribute Pattern | MongoDB**  
<https://www.mongodb.com/blog/post/building-with-patterns-the-attribute-pattern> ▾  
Jan 24, 2019 - Learn about the Attribute Schema Design pattern in MongoDB. This pattern is used to target similar fields in a document and reducing the ...
- Building with Patterns: The Bucket Pattern | MongoDB**  
<https://www.mongodb.com/blog/post/building-with-patterns-the-bucket-pattern> ▾  
Jan 31, 2019 - Learn about the Bucket Schema Design pattern in MongoDB. This pattern is used to group similar documents together for easier aggregation.
- Building With Patterns: The Outlier Pattern | MongoDB**  
<https://www.mongodb.com/blog/post/building-with-patterns-the-outlier-pattern> ▾  
3 days ago - So far in this Building with Patterns series, we've looked at the Polymorphic, Attribute, and Bucket patterns. While the document schema in ...



# MongoDB Atlas

11:00 – 12:00

---



# MongoDB 101 Workshop

---

## Introduction to MongoDB (30 mins)

### Setting up MongoDB (30 mins)

- Laptop: Standalone Engine & Shell, Compass, MongoDB Driver (Python)
- Data Generator: MGenerateJS
- Cloud: MongoDB Atlas account

### Querying, Indexing and Aggregation (30 mins)

- Using Compass, MongoDB Shell & Python

### Schema Design & Use Cases (30 mins)

- Embedding & Referencing, 1-1, 1-M, M-M, Patterns
- Single View, IoT, Time Series, eCommerce/Content Management

### MongoDB Atlas (60 mins)

- Management, Clusters, Security
- High Availability & Scalability ( Replication & Sharding )
- Managed BI Connector & Charts
- ODBC for BI Connector
- MongoDB Stitch\*



# Atlas

unlocks **agility**  
and **reduces**  
**cost**

---



Self-service and  
elastic



Global and highly  
available



Secure by default



Comprehensive  
monitoring



Managed backup



Cloud agnostic



# MongoDB Atlas: Database as a Service

## Self-service and elastic

- Deploy in minutes
- Scale up/down without downtime
- Automated upgrades

## Global and highly available

- 59 Regions worldwide
- Replica sets optimized for availability
- Global clusters for read/write everywhere deployments

## Secure by default

- Network isolation and Peering
- Encryption in flight and at rest with key management
- Role-based access control; auditing
- SOC 2 / Privacy Shield / HIPAA

## Comprehensive Monitoring

- Performance Advisor
- Dashboards w/ 100+ metrics
- Real Time Performance
- Customizable alerting

## Managed Backup

- Point in Time Restore
- Queryable backups
- Consistent snapshots

## Cloud Agnostic

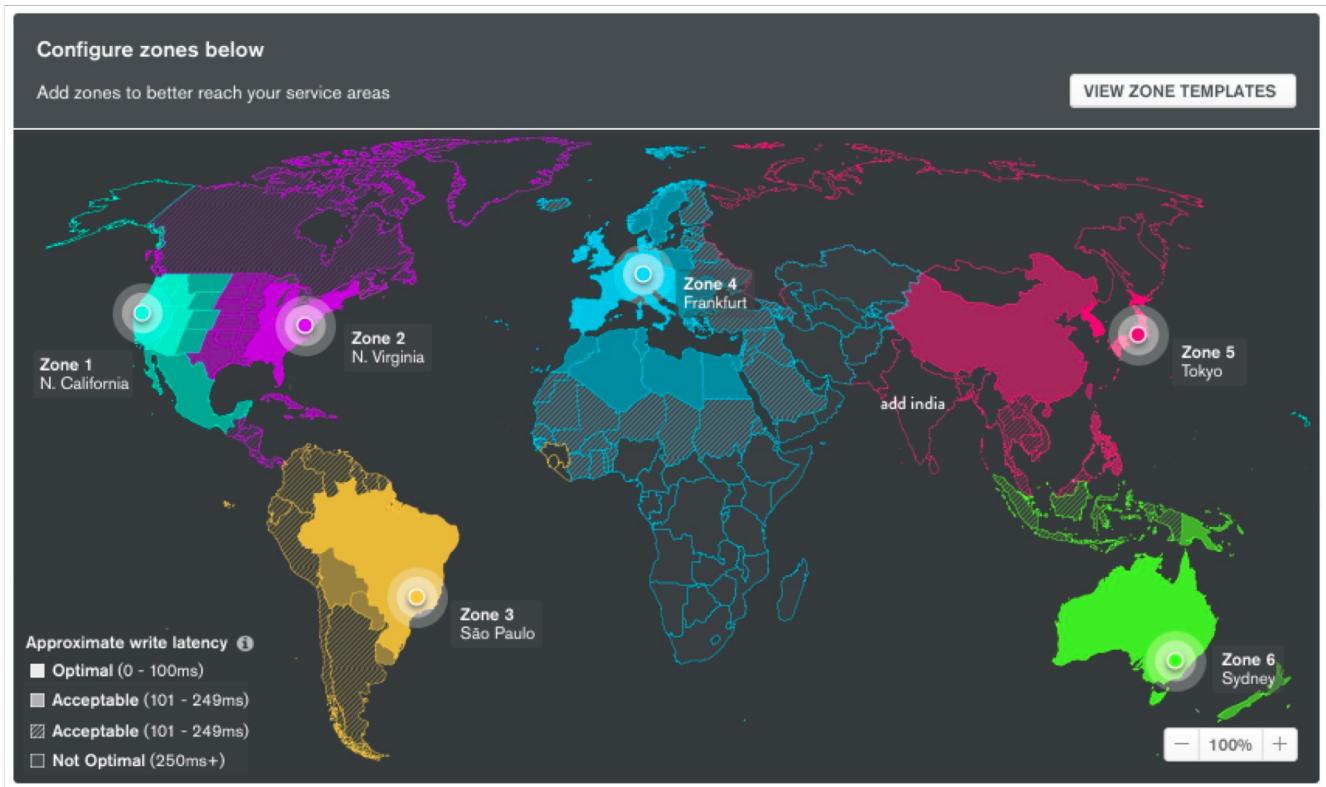
- AWS, Azure, and GCP
- Easy migrations
- Consistent experience

# What's New: Atlas Global Clusters

---

Distribute your fully automated database across multiple geographically distributed zones made up of one or more cloud regions

- Read and write locally to provide single-digit millisecond latency for your distributed applications
- Ensure that data from certain geographies lives in predefined zones
- Easily deploy using prebuilt zone templates or build your own zones by choosing cloud regions in an easy-to-use, visual interface



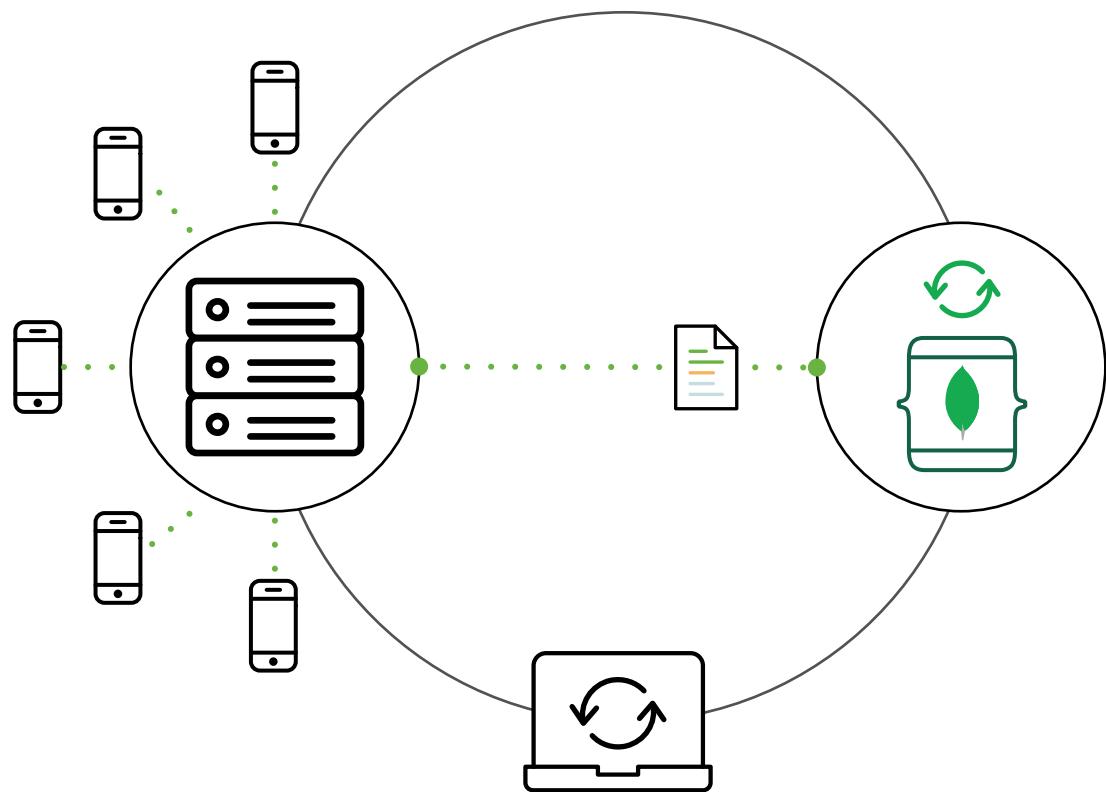
# What's New: MongoDB Atlas Enterprise Security

---



- LDAP integration
- Database level auditing
- Encrypted storage engine: BYO Key Management Service (AWS initially)
- HIPAA Eligible

# Freedom to run Anywhere: Cloud Migration



Migrate existing deployments running anywhere into MongoDB Atlas with minimal impact to your application.

Live migration works by:

- Performing a sync between your source database and a target database hosted in MongoDB Atlas
- Syncing live data between your source database and the target database by tailing the oplog
- Notifying you when it's time to cut over to the MongoDB Atlas cluster



# mongoDB® Atlas : The only *true* multi-cloud database as a service



Google Cloud Platform

## Database that runs the same everywhere

Consistent experience  
across AWS, Azure, and  
GCP

## Coverage in any geography

Deploy in 59 regions  
worldwide

Create globally distributed  
databases with a few clicks

## Leverage the benefits of a multi-cloud strategy

Exploit the benefits of AWS,  
Azure, or GCP services on  
your data

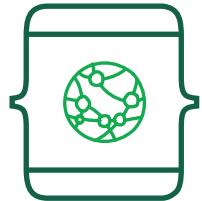
## Avoid lock-in

Easily migrate data  
between cloud providers



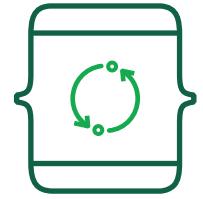
# Intelligently put data where you need it

---



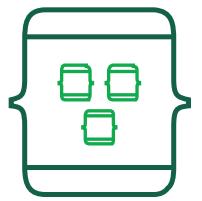
## Highly Available

Built-in **multi-region** high availability, replication & automated failover



## Workload Isolation

Ability to run both **operational & analytics workloads** on same cluster, for timely insight and lower cost



## Scalability

Elastic horizontal scalability – add/remove capacity dynamically **without downtime**

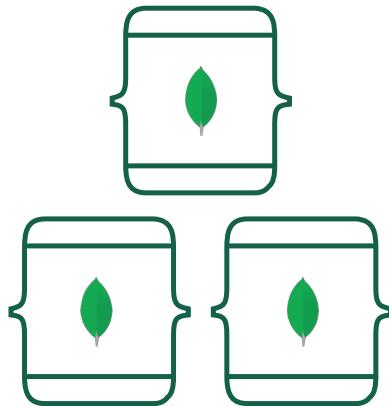


## Locality

Declare **data locality rules** for governance (e.g. data sovereignty), class of service & local low latency access

# Put data where you need it: Availability

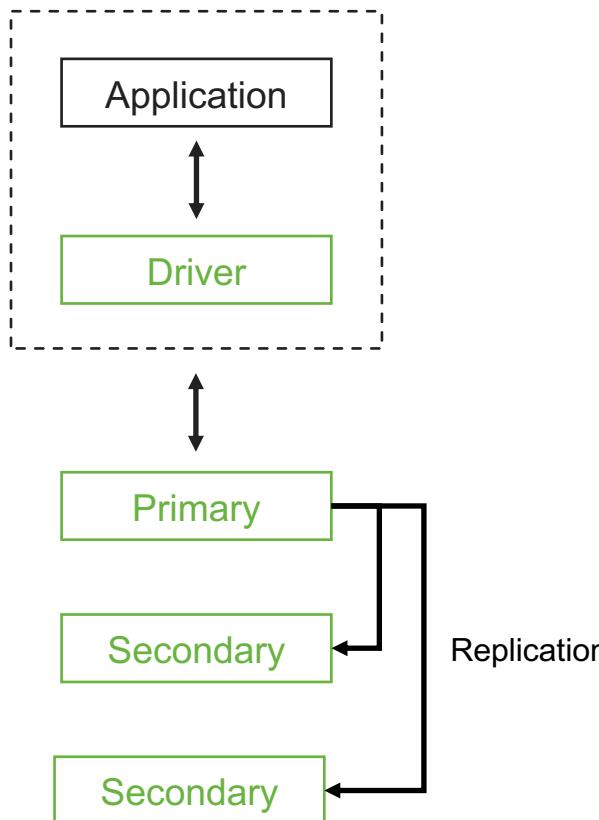
---



## Replica Sets

- Up to 50 replicas, distributed across racks, data centers, and regions
- Self-healing using RAFT-based consensus protocol for automated failover & recovery
- Tunable durability and consistency controls
- Always-on write availability with retryable writes

# MongoDB Replica Set Architecture



**Replica Set – 2 to 50 copies**

**Self-healing**

**Data Center Aware**

**Addresses availability considerations:**

- High Availability
- Disaster Recovery
- Maintenance

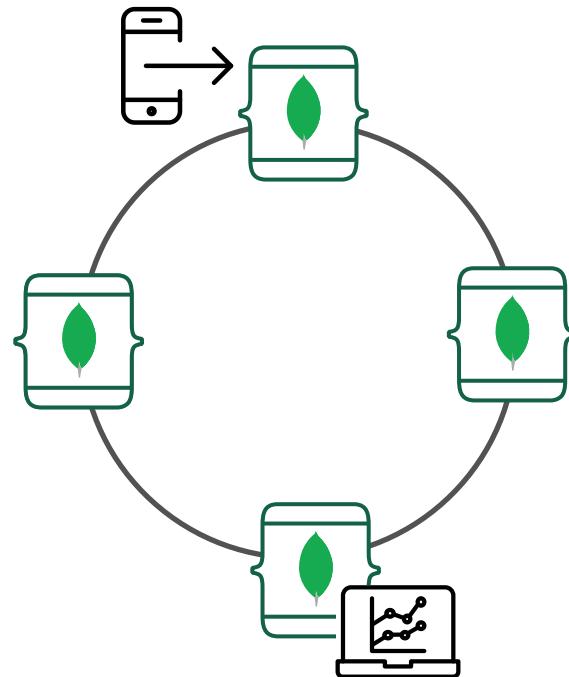
**Workload Isolation: operational & analytics**

# MongoDB's native replica sets puts your entire database right next to users



# Put data where you need it: Workload Isolation

---

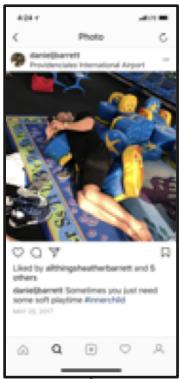


## Enable different workloads on the same data

- Combine operational and analytical workloads on a single data platform
- Extract live insights from real-time data to enrich applications
- One set of nodes serving operational apps, replicating to dedicated nodes serving analytics: up to 50 nodes in a single replica set
- ETL-free

# Co-locating operational and analytical workloads

## Transactional



Primary

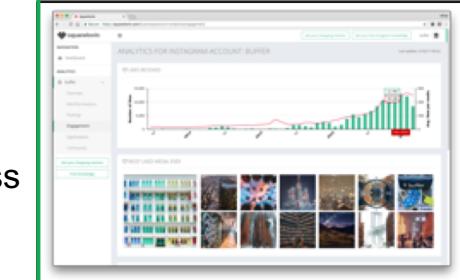
Secondary

Secondary

Secondary  
{use = analytics}

Secondary  
{use = analytics}

BI & Reporting  
mongoDB Charts  
mongoDB Compass  

BI Connector

Predictive Analytics & Data Science



Aggregations



python™



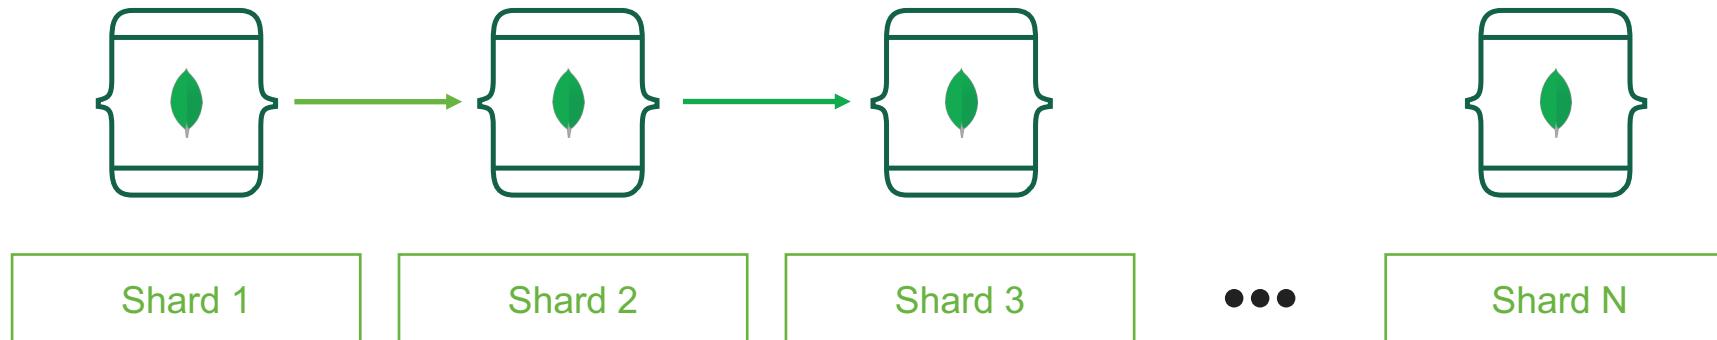
# Technical Implementation

---

```
// Set up replica set
rs.initiate(
{
  "_id": "rsl",
  "version": 1,
  "members":
  [
    {"priority": 1, "host": "localhost:27001", "_id": 0, "tags": {"use": "op"}},
    {"priority": 1, "host": "localhost:27002", "_id": 1, "tags": {"use": "op"}},
    {"priority": 1, "host": "localhost:27003", "_id": 2, "tags": {"use": "op"}},
    {"priority": 0, "host": "localhost:27004", "_id": 3, "tags": {"use": "analytics"}},
    {"priority": 0, "host": "localhost:27005", "_id": 4, "tags": {"use": "analytics"}}
  ]
})

// Force reads to analytics servers
read_client = MongoClient(host = 'localhost:27001', replicaset = 'rsl',
  readPreference='secondary', readPreferenceTags = ["use:analytics"])
```

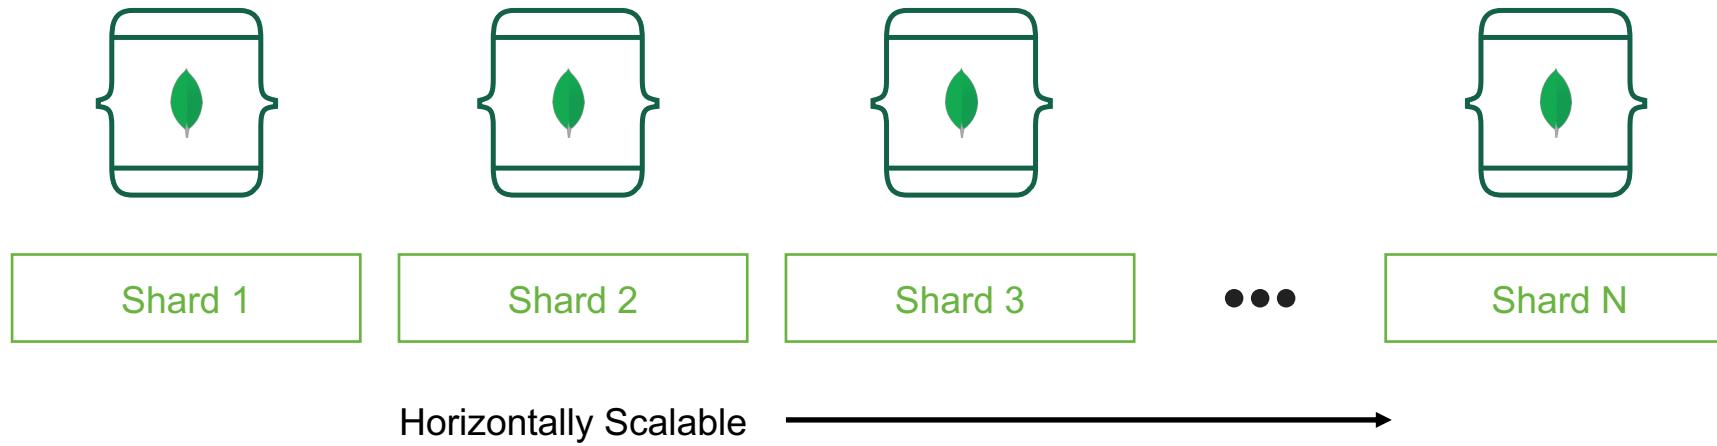
# Put data where you need it: Scalability



## Auto-Sharding

- Automatically scale beyond the constraints of a single node
- Application transparent
- Scale and rebalance incrementally, in real time
- Unlike NoSQL systems that randomly spray data across a cluster, MongoDB exposes multiple data distribution policies to optimize for query patterns and locality

# Scaling MongoDB: Automatic Sharding

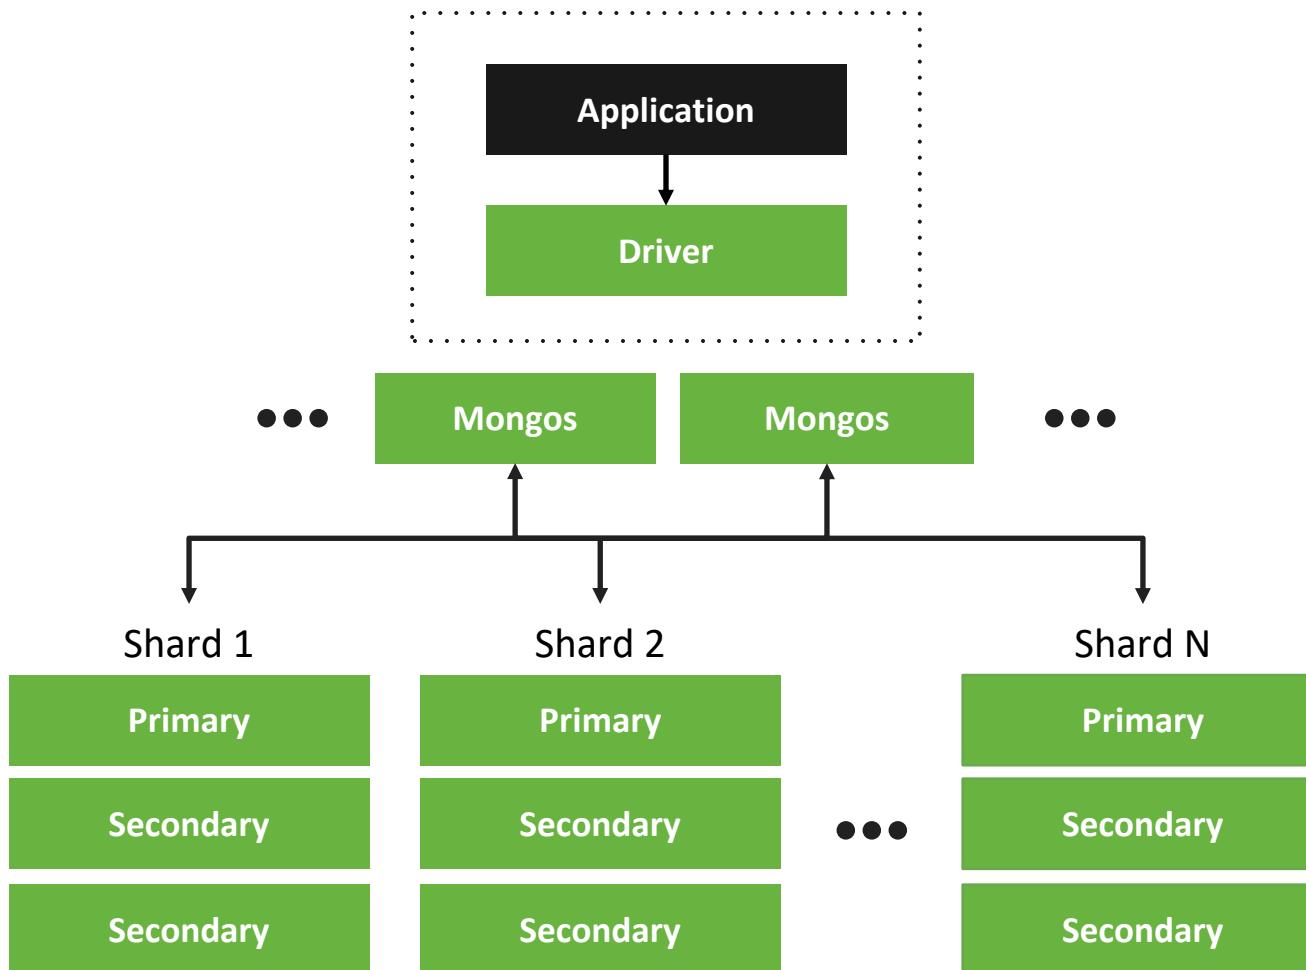


Multiple sharding policies: hashed, ranged, zoned

Increase or decrease capacity as you go

Automatic balancing for elasticity

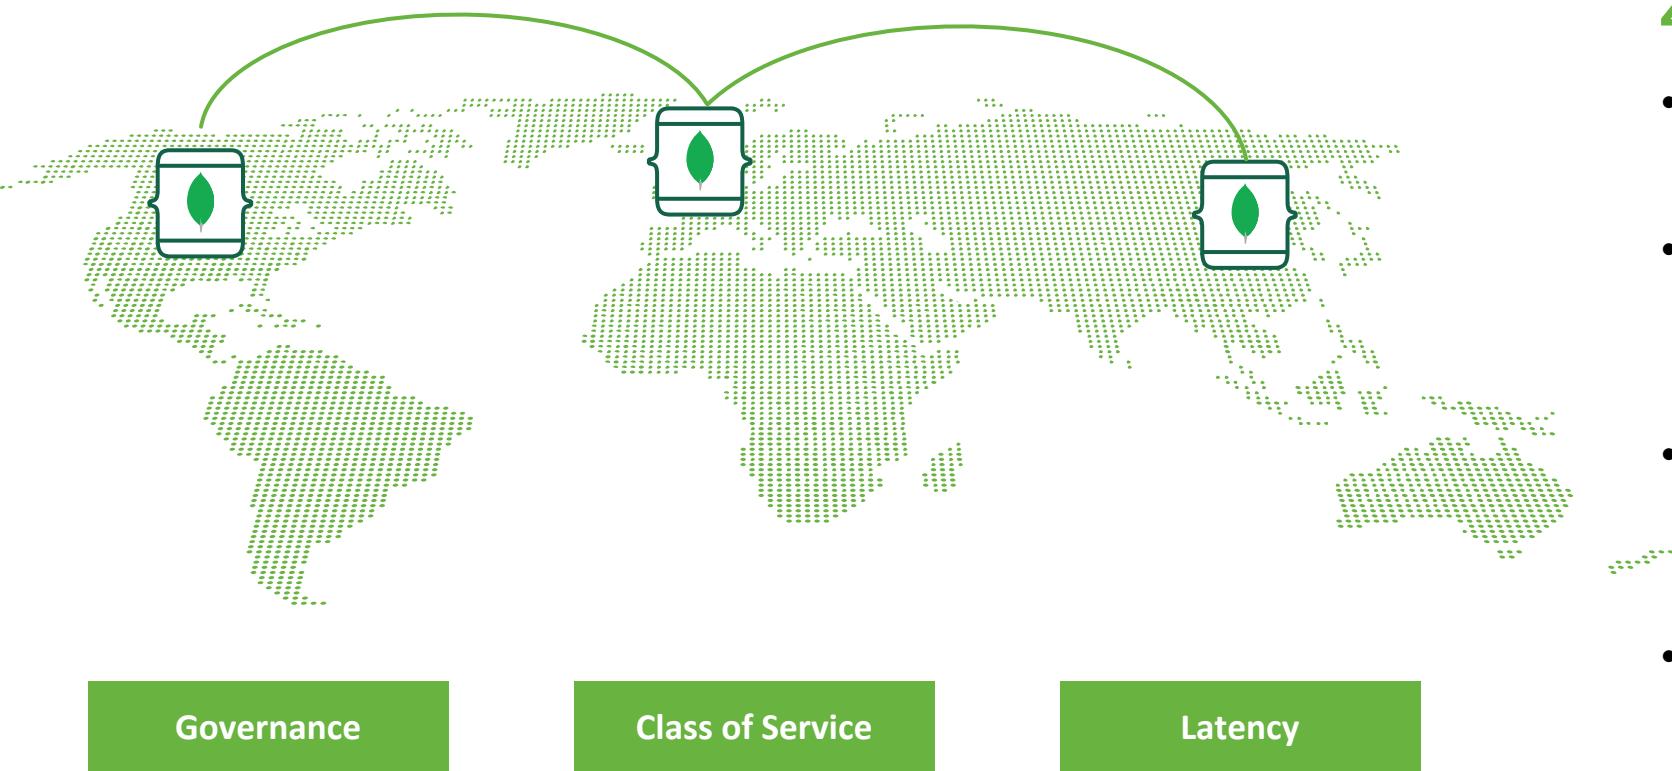
# Sharding Architecture



**High availability**  
- Replica sets

**Horizontal scalability**  
- Sharding

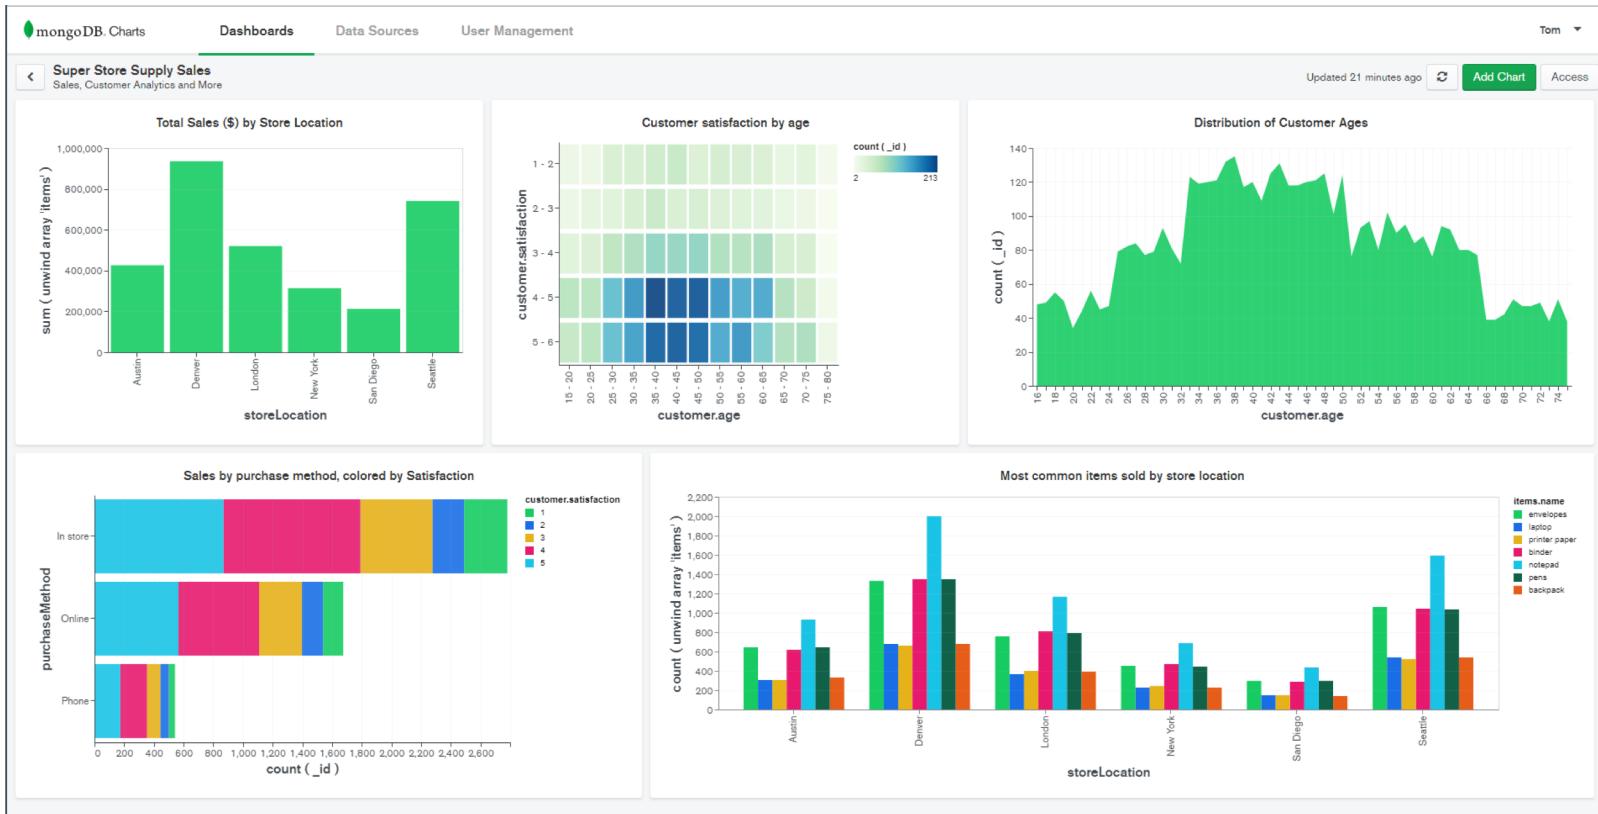
# Put data where you need it: Locality



## Intelligent Distribution via Zoned Sharding

- Policies to define data placement
- Name a server by region, tag your documents by region and MongoDB does the rest
- Documents automatically migrated as shard key ranges are modified
- Global queries across all data in all zones

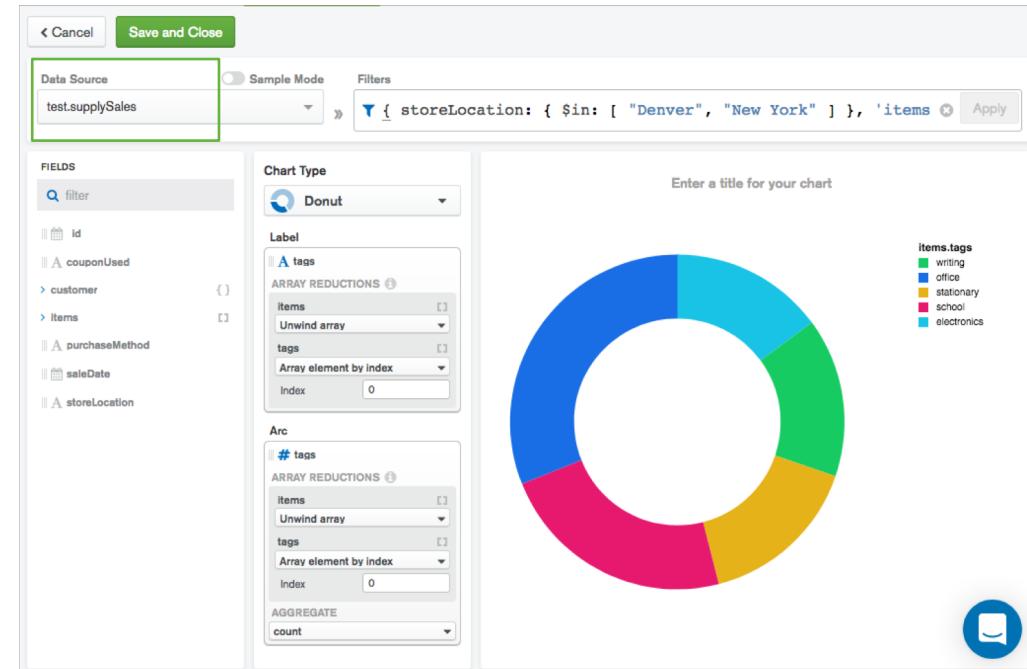
# Sophisticated Analytics & Visualizations Of Data In Place



- Rich MongoDB query language & idiomatic drivers
- Connector for BI
- Connector for Spark
- Charts (beta)

# MongoDB Charts: Create, Visualize, Share

```
{  
  "id": {"$oid": "5afbc3dc09c8d2dd5852cf2"},  
  "saleDate": {"$date": "2017-11-08T19:06:53.449Z"},  
  "items": [  
    {  
      "name": "envelopes",  
      "tags": ["stationery", "office", "general"],  
      "price": {"$numberDecimal": "9.83"},  
      "quantity": 10  
    },  
    {  
      "name": "pens",  
      "tags": ["office", "writing", "school", "stationary"],  
      "price": {"$numberDecimal": "73.62"},  
      "quantity": 2  
    },  
    {  
      "name": "laptop",  
      "tags": ["office", "school", "electronics"],  
      "price": {"$numberDecimal": "595.72"},  
      "quantity": 4  
    },  
    {  
      "name": "notepad",  
      "tags": ["office", "writing", "school"],  
      "price": {"$numberDecimal": "34.65"},  
      "quantity": 3  
    }  
  ],  
  "storeLocation": "Seattle",  
  "customer": {  
    "gender": "M",  
    "age": 45,  
    "email": "uga@we.soo",  
    "satisfaction": 4  
  },  
  "couponUsed": false,  
  "purchaseMethod": "Online"  
}
```



Work with complex data

Connect to data sources securely.  
Filter. Sample. Visualize.

Share dashboards and  
collaborate

# MongoDB Local Replica Set Workshop

- **Replica-set.txt**

```
mongod
2019-02-18T22:14:48.459-0700 I NETWORK [conn14] Skip closing connection for connection # 7
2019-02-18T22:14:48.459-0700 I NETWORK [conn14] Skip closing connection for connection # 5
2019-02-18T22:14:48.459-0700 I NETWORK [conn14] Skip closing connection for connection # 4
2019-02-18T22:14:48.459-0700 I NETWORK [conn14] Skip closing connection for connection # 2
2019-02-18T22:14:48.459-0700 I REPL [conn14] Handing off election to localhost:27001
2019-02-18T22:14:48.459-0700 I NETWORK [conn14] Error sending response to client: SocketException: Broken pipe. Ending connection from 127.0.0.1:51441 (connection id: 14)
2019-02-18T22:14:48.459-0700 I NETWORK [conn14] end connection 127.0.0.1:51441 (6 connections now open)
2019-02-18T22:14:48.462-0700 I NETWORK [listener] connection accepted from 127.0.0.1:51442 #15 (7 connections now open)
2019-02-18T22:14:48.462-0700 I NETWORK [conn15] received client metadata from 127.0.0.1:51442 conn15: { application: { name: "MongoDB Shell" }, driver: { name: "MongoDB Internal Client", version: "4.0.6" }, os: { type: "Darwin", name: "Mac OS X", architecture: "x86_64", version: "18.2.0" } }
2019-02-18T22:14:48.712-0700 I REPL [replexec-1] Member localhost:27001 is now in state PRIMARY
2019-02-18T22:14:50.530-0700 I REPL [rsBackgroundSync] sync source candidate: localhost:27001
2019-02-18T22:14:50.531-0700 I ASIO [RS] Connecting to localhost:27001
2019-02-18T22:14:50.533-0700 I REPL [rsBackgroundSync] Changed sync source from empty to localhost:27001
2019-02-18T22:14:50.534-0700 I ASIO [RS] Connecting to localhost:27001
2019-02-18T22:14:57.490-0700 I NETWORK [conn4] end connection 127.0.0.1:51378 (6 connections now open)

1. mongo
2019-02-18T22:14:13.336-0700 I REPL [rsBackgroundSync] sync source candidate: localhost:27000
2019-02-18T22:14:13.336-0700 I ASIO [RS] Connecting to localhost:27000
2019-02-18T22:14:13.338-0700 I REPL [rsBackgroundSync] Changed sync source from empty to localhost:27000
2019-02-18T22:14:13.349-0700 I ASIO [RS] Connecting to localhost:27000
2019-02-18T22:14:49.377-0700 I REPL [replexec-0] Member localhost:27000 is now in state SECONDARY
2019-02-18T22:14:49.377-0700 I REPL [replexec-1] Member localhost:27001 is now in state PRIMARY
2019-02-18T22:14:50.536-0700 I REPL [replication-1] Restarting oplog query due to error: InterruptedDueToReplStateChange: error in fetcher batch callback :: caused by :: operation was interrupted. Last fetched optime (with hash): { ts: Timestamp(1550553282, 1), t: 4 } [-6117836752632269864]. Restart s remaining: 1
2019-02-18T22:14:50.536-0700 I REPL [replication-1] Scheduled new oplog query Fetcher source: localhost:27000 database: local query: { find: "oplog.rs", filter: { ts: { $gte: Timestamp(1550553282, 1) } }, tailable: true, oplogReplay: true, awaitData: true, maxTimeMS: 2000, batchSize: 13981010, term: 5, readConcern: { afterClusterTime: Timestamp(1550553282, 1) } } query metadata: { $repData: 1, $oplogQueryData: 1, $readPreference: { mode: "secondaryPreferred" } } active: 1 findNetworkTimeout: 7000ms getMoreNetworkTimeout: 10000ms shutting down?: 0 first: 1 firstCommandScheduler: RemoteCommandRetryScheduler request: RemoteCommand { target:localhost:27000 db:local cmd: { find: "oplog.rs", filter: { ts: { $gte: Timestamp(1550553282, 1) } }, tailable: true, oplogReplay: true, awaitData: true, maxTimeMS: 2000, batchSize: 13981010, term: 5, readConcern: { afterClusterTime: Timestamp(1550553282, 1) } } } active: 1 callbackHandle.valid: 1 callbackHandle.cancelled: 0 attempt: 1 retryPolicy: RetryPolicyImpl maxAttempts: 1 maxTimeMillis: -1ms
[{"configVersion": 1, "ok": 1, "operationTime": Timestamp(1550553282, 1), "$clusterTime": {"clusterTime": Timestamp(1550553282, 1), "signature": {"hash": BinData(0, "AAAAAAAAAAAAAAAAAAAAAAA="), "keyId": NumberLong(0)}}, "rs0:PRIMARY> rs.stepDown()"}, {"2019-02-18T22:14:48.460-0700 E QUERY [js] Error: error doing query: failed: network error while attempting to run command 'replSetStepDown' on host '127.0.0.1:27000' : DB.prototype.runCommand@src/mongo/shell/db.js:168:1 DB.prototype.adminCommand@src/mongo/shell/db.js:186:1 rs.stepDown@src/mongo/shell/utils.js:1444:12 @shell):1:1", "2019-02-18T22:14:48.462-0700 I NETWORK [js] trying reconnect to 127.0.0.1:27000 failed", "2019-02-18T22:14:48.462-0700 I NETWORK [js] reconnect 127.0.0.1:27000 ok", "rs0:SECONDARY>"]
```



# MongoDB Atlas Workshop

---

- **Setup M0 Cluster**
- Insert data: Customer Single View Data Generator / Mongolimport / IP Whitelist
- Enable Charts
- Build a dashboard
- Repeat previous exercises against Atlas (shell or python)

Optional (credit card required)

- **Scale-up to M10**
- Enable BI Connector
- Setup ODBC Connector
- Create dashboard using Tableau, QlikView, etc.



mongoDB®

FOR GIANT IDEAS